



CS109: Probability for Computer Scientists

Chris Piech

My parents are interesting folks

I originally concentrated in graphics and worked at Pixar

- Childhood: Nairobi, Kenya
- High School: Kuala Lumpur, Malaysia
- Stanford University Ph.D. in Deep Learning
- Research lab on AI for Social Good
- Researcher in the MSR cognitive science lab

The problem I really want to solve is to make high quality more education accessible



I Took the First CS109 Class



Back when I looked like this 😊



Teaching Assistants

Several outstanding TAs:

William Monroe (Head TA)

William Chen

Govinda Dasu

Luke Johnston

Winnie Lin

Michela Meister

Priyanka Nigam

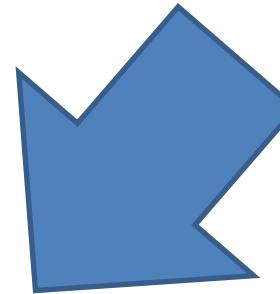
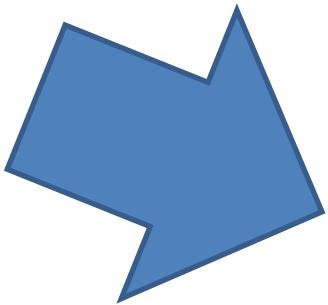
Benjamin Ulmer



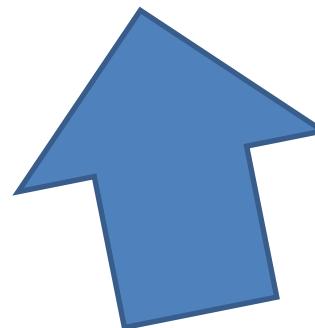
Who are you?

Course mechanics
(this is a light version. Please read the handout
for details).

Essential Information



cs109.stanford.edu



Are you in the right place?

Prereqs

What you really need:

CS106B/X:

- Recursion
- Hash Tables
- Binary Trees
- Programming

CS103:

- Proof techniques (induction)
- Set theory

Math 51 or CME 100

- Multivariate differentiation
- Multivariate integration
- Basic facility with linear algebra (vectors)

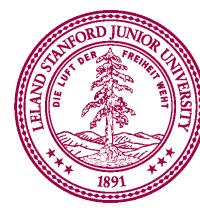
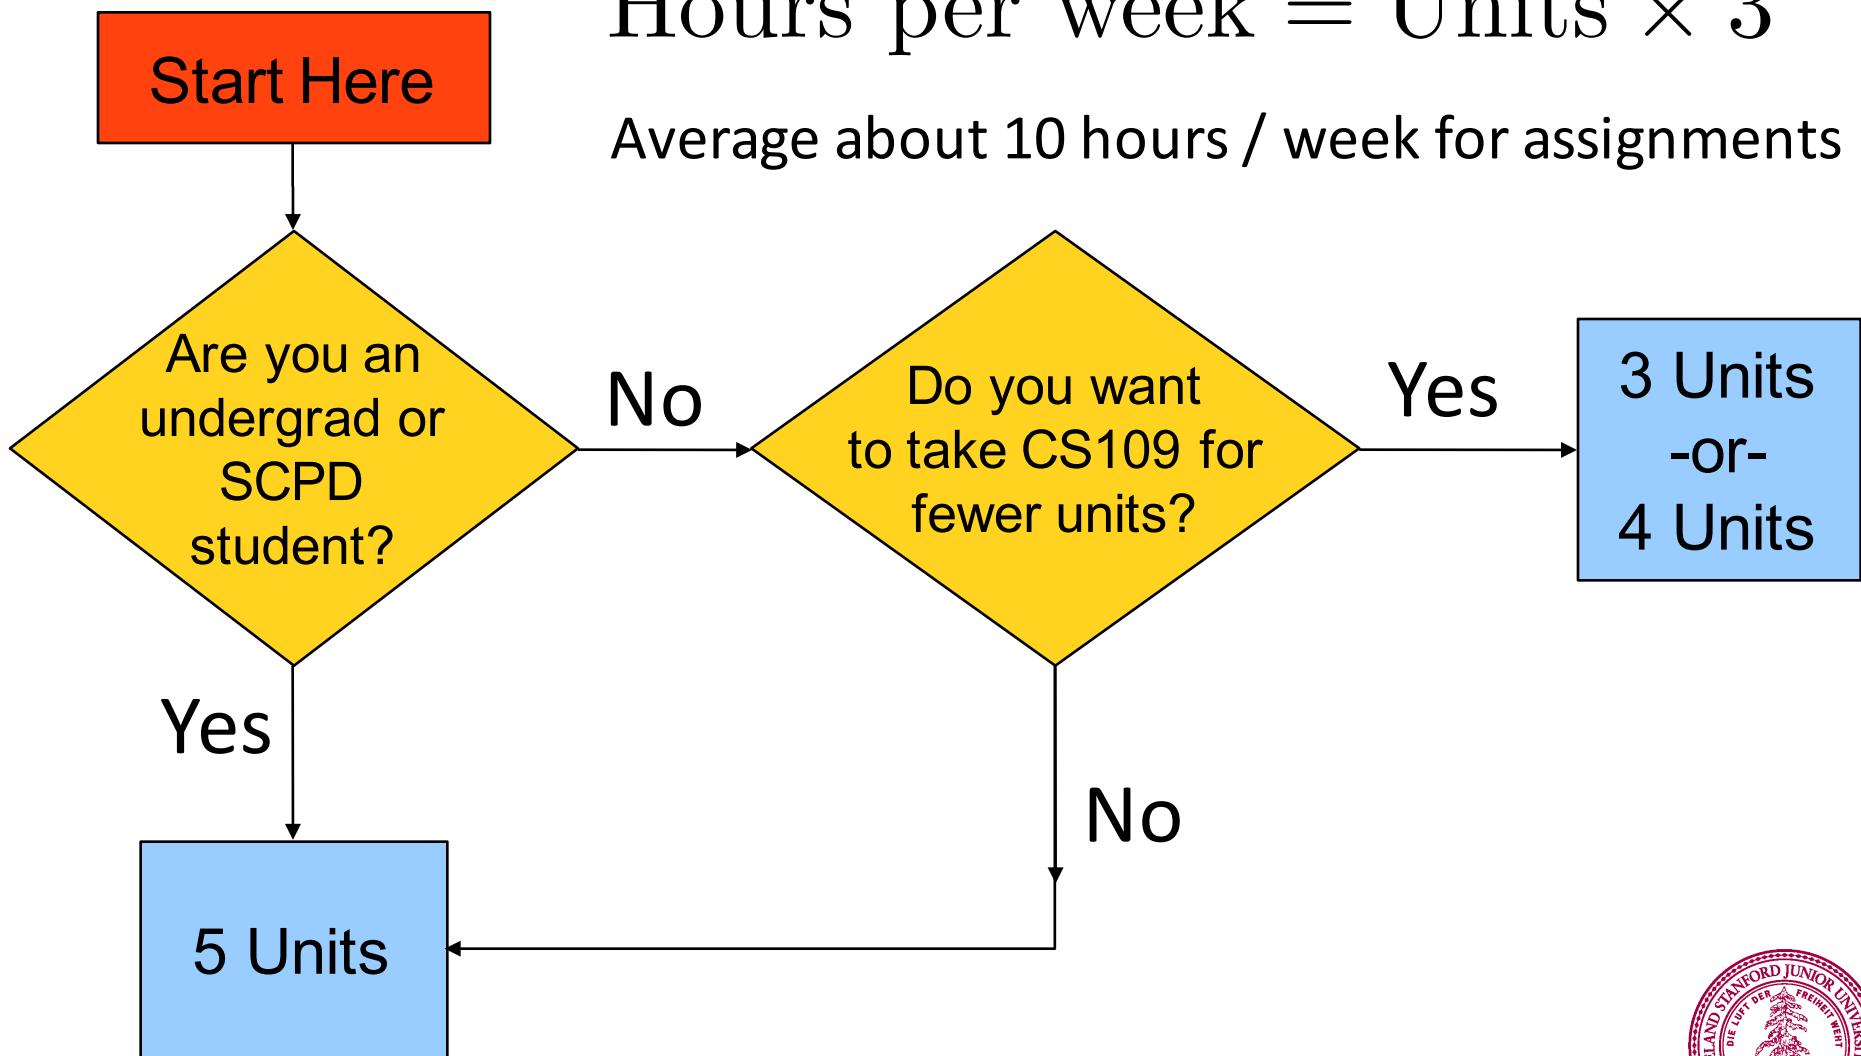


Staff Contact

- Post to Piazza
- Go to Working Office Hours
- Email cs109@cs.stanford.edu
- Email Chris or go to his office for course level issues.



CS109 Units



Video Taped



Piech, CS106A, Stanford University



Class Breakdown

45%

6 Assignments

20%

Midterm

Tuesday May 9th, 7-9pm

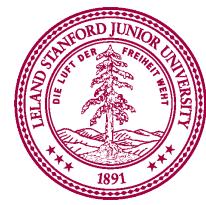
35%

Final

Tuesday June 13th, 8:30-11:30am

5%

Lecture Participation



Late Days

2



The Student Honor Code

Story of Modern AI

Modern AI
or, How we learned to combine
probability and programming

Brief History



Narrow Intelligence

Play Chess

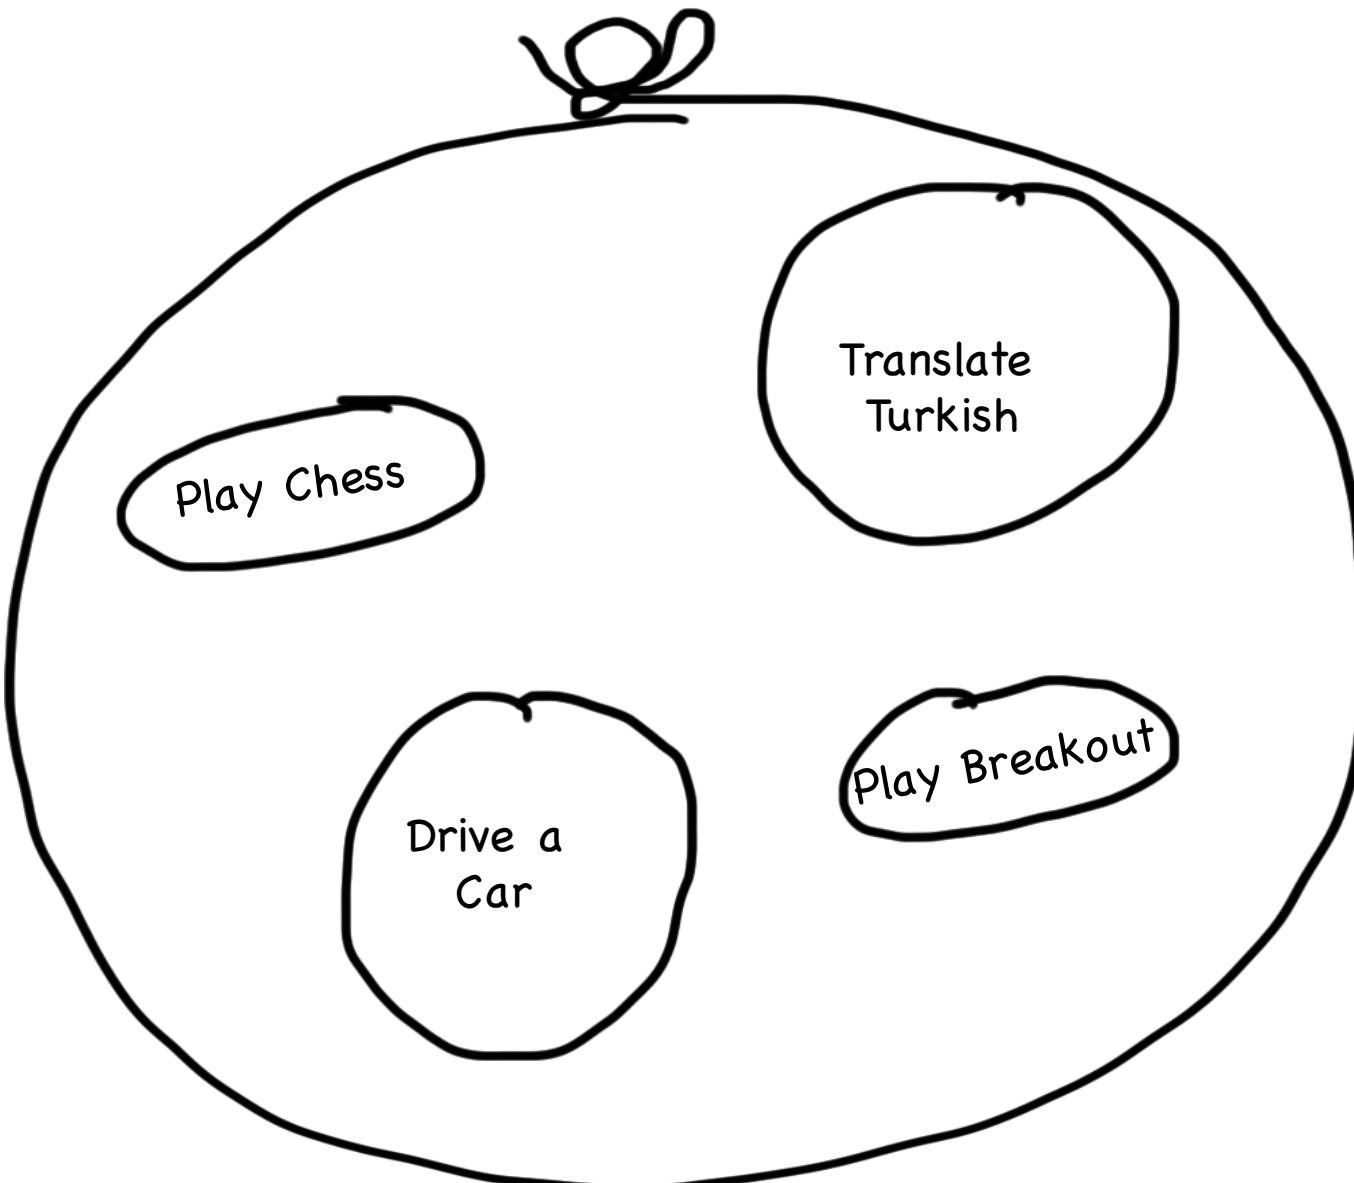
Translate
Turkish

Drive a
Car

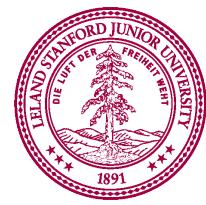
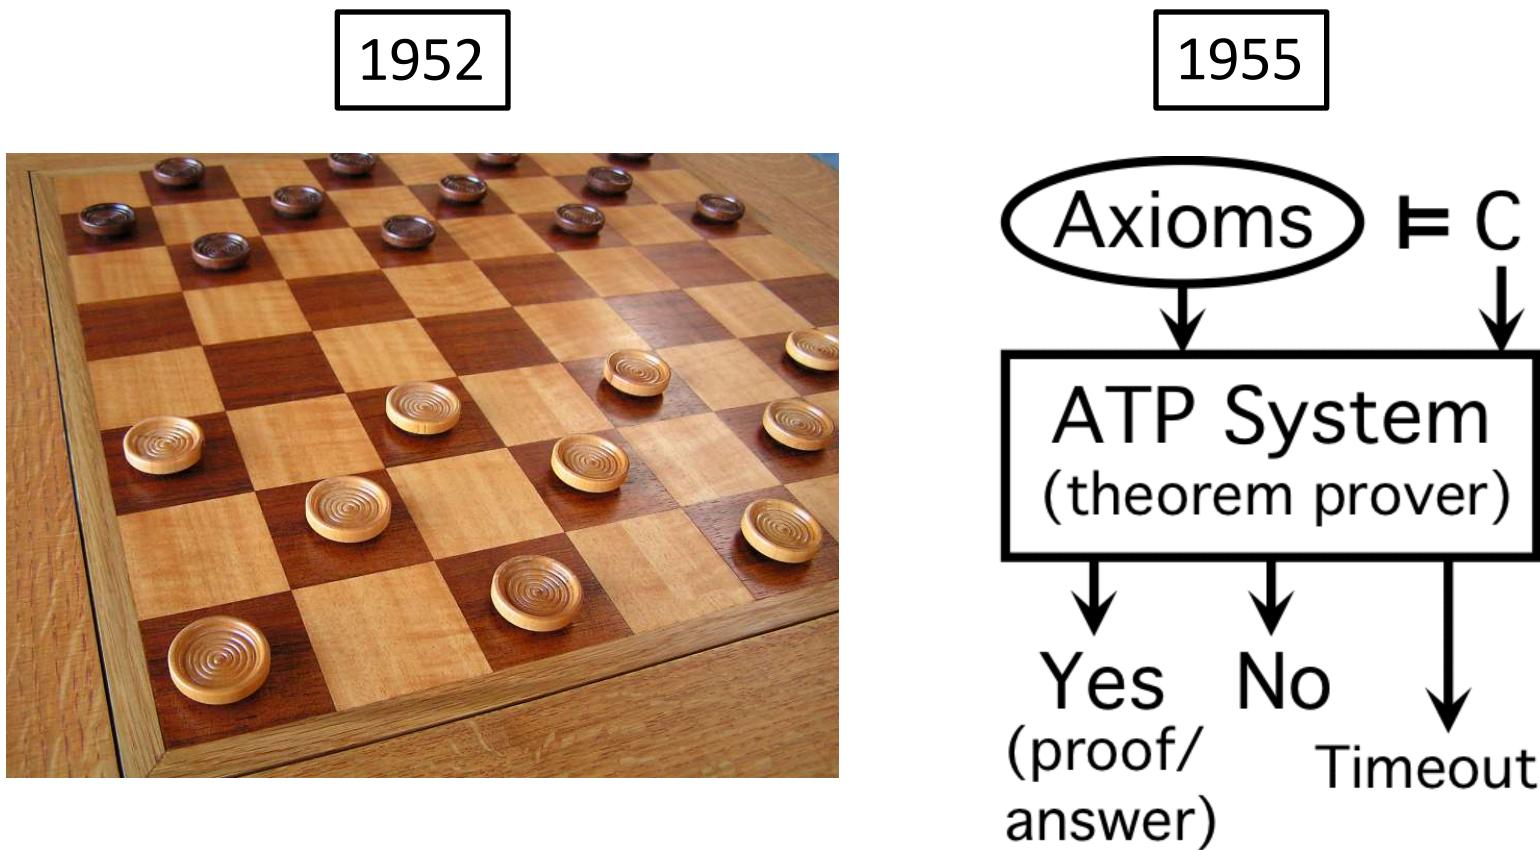
Play Breakout



General Intelligence



Early Optimism 1950



Early Optimism 1950

“Machines will be capable,
within twenty years, of doing
any work a man can do.”
–Herbert Simon, 1952



Underwhelming Results 1950s to 1980s

The spirit is willing but the flesh is weak.



(Russian)



The vodka is good but the meat is rotten.

The world is too complex



BRACE YOURSELVES



WINTER IS COMING

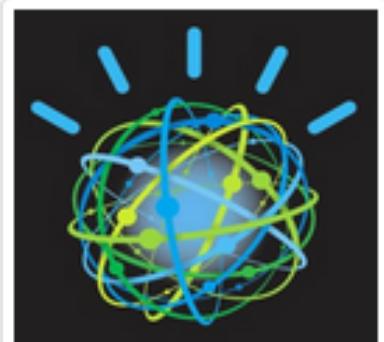
Big Milestones Pt 1



1997 Deep Blue

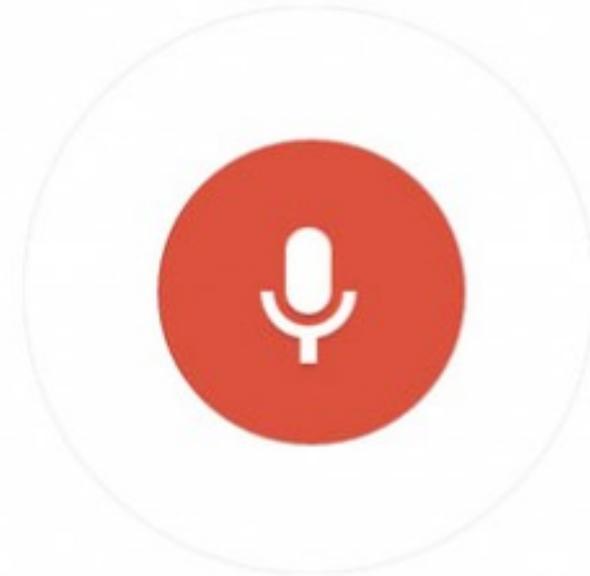


2005 Stanley



2011 Watson

Big Milestones Pt 2



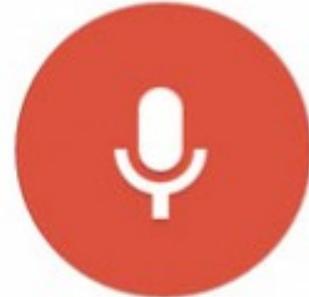
Piech, CS106A, Stanford University



Told Vision Was 30 Years Out



Told Speech Was 30 Years Out



Almost perfect...

Piech, CS106A, Stanford University



Huge Progress



Passed the Turing Test?



EUGENE GOOSTMAN THE WEIRDEST CREATURE IN THE WORLD

33% of judges thought Eugene was a human

When only given 5 minutes...

Piech, CS106A, Stanford University



What is going on?

Two Great Ideas

1. Probability from Examples

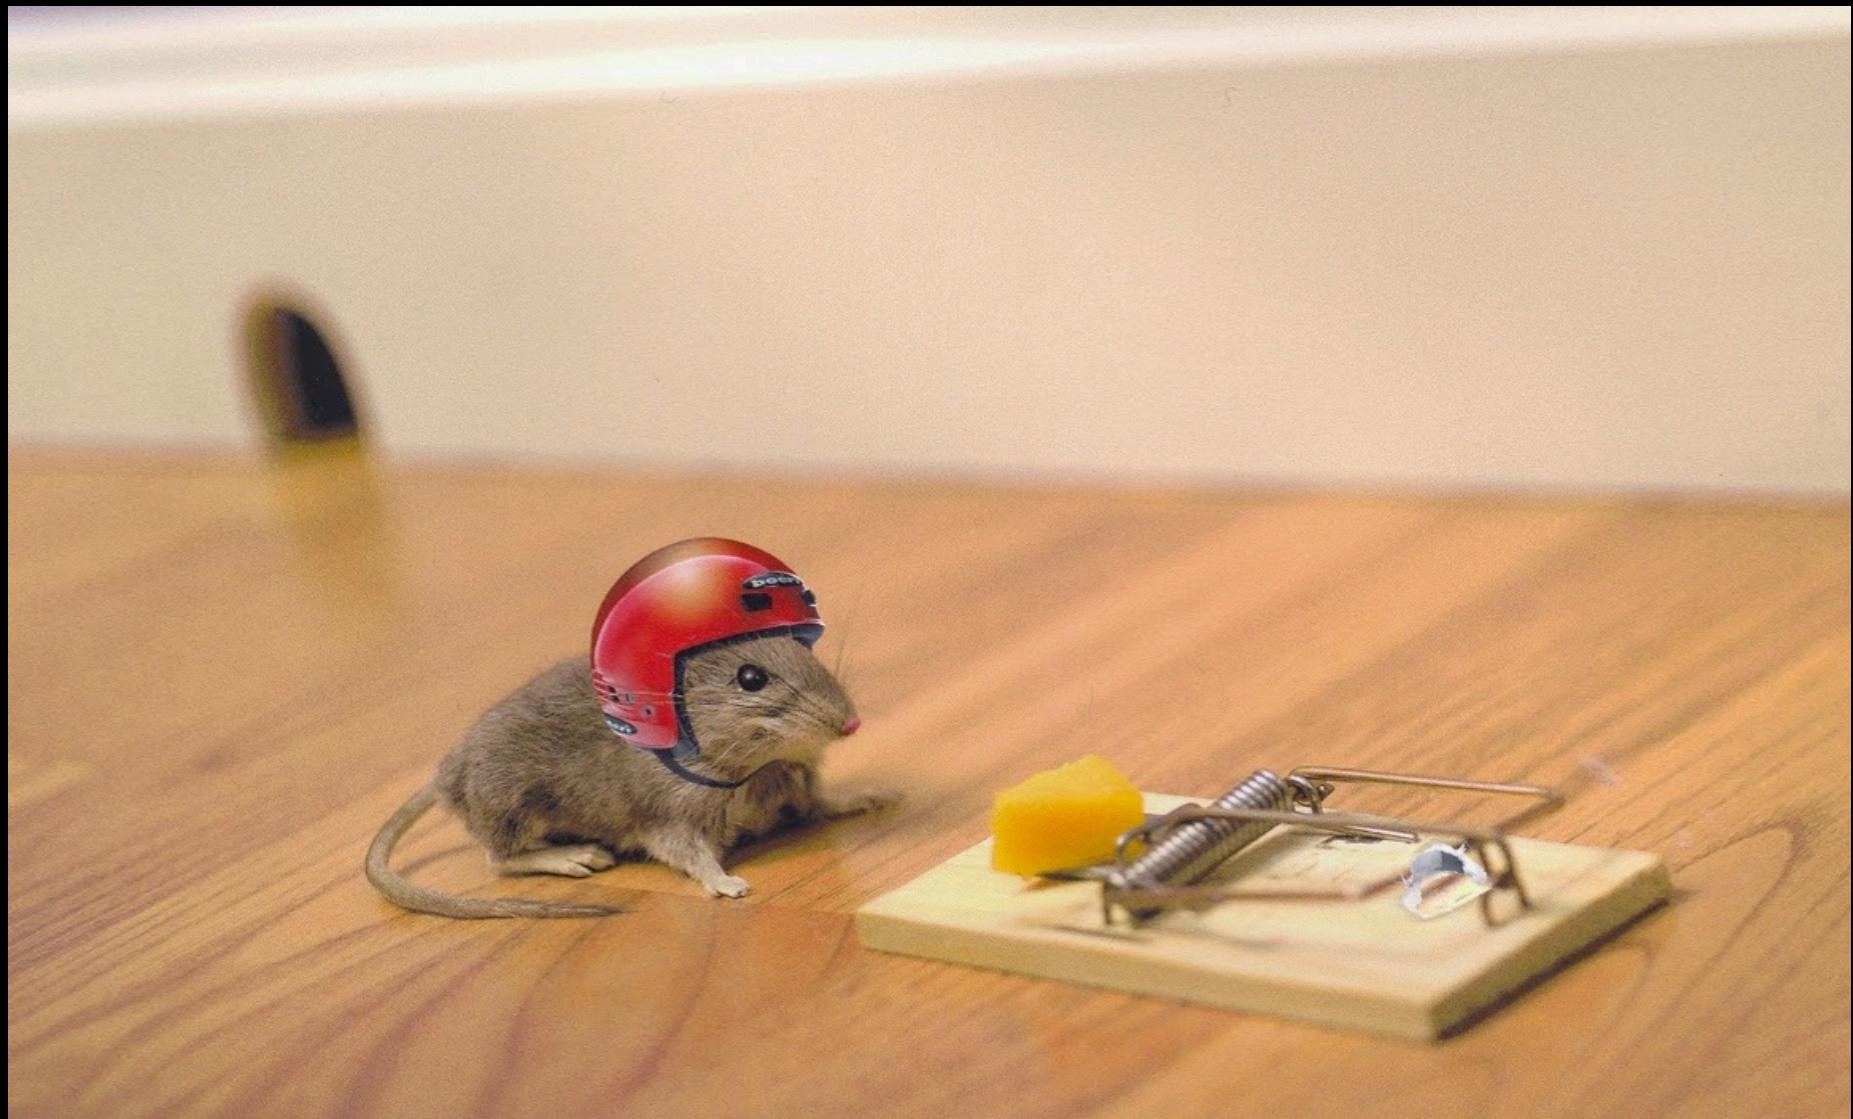
2. Artificial Neurons

Two Great Ideas

1. Probability from Examples

2. Artificial Neurons

1. Probability From Examples



When Does the Magic Happen?

Lots of
Data + Sound
Probability

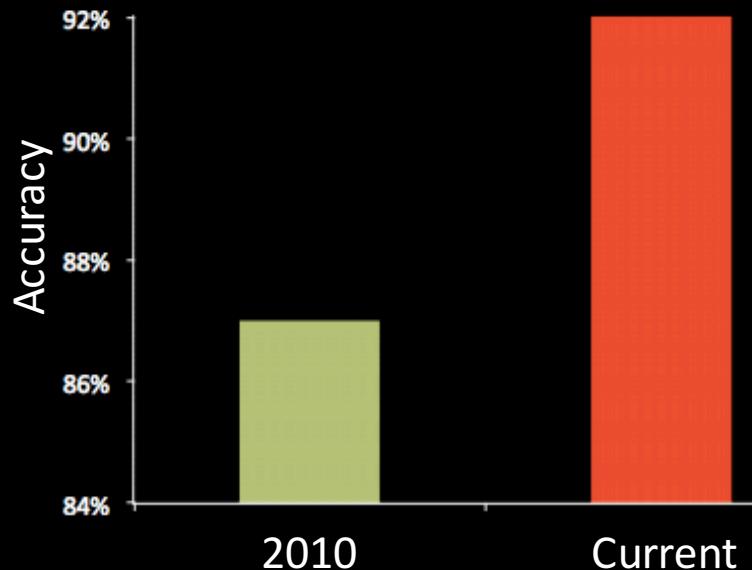
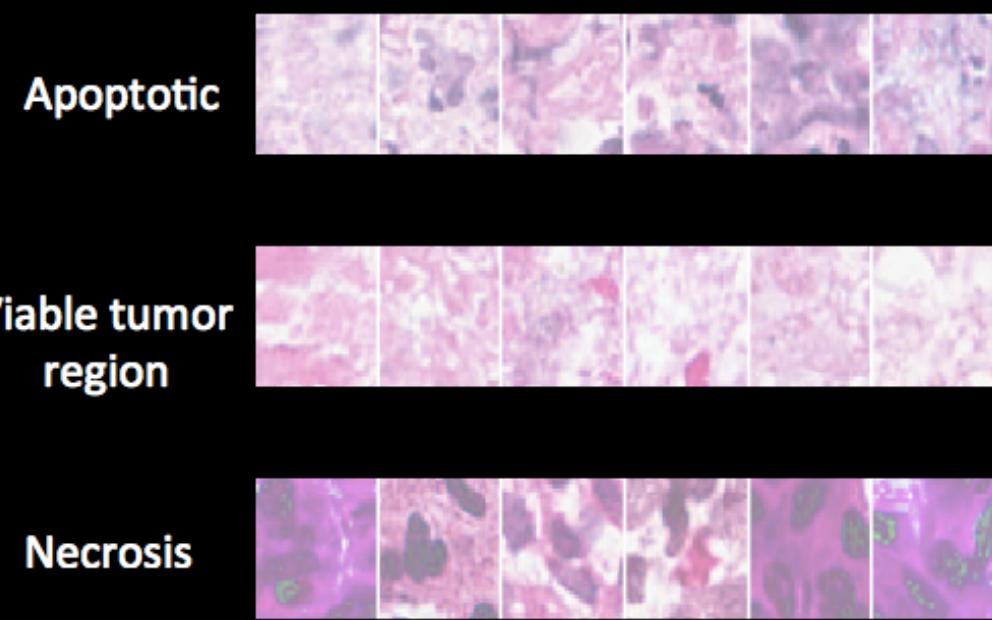


Machine Learning

Basically just a rebranding of statistics
and probability.

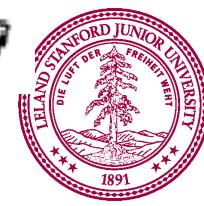


Vision: Detect Cancer

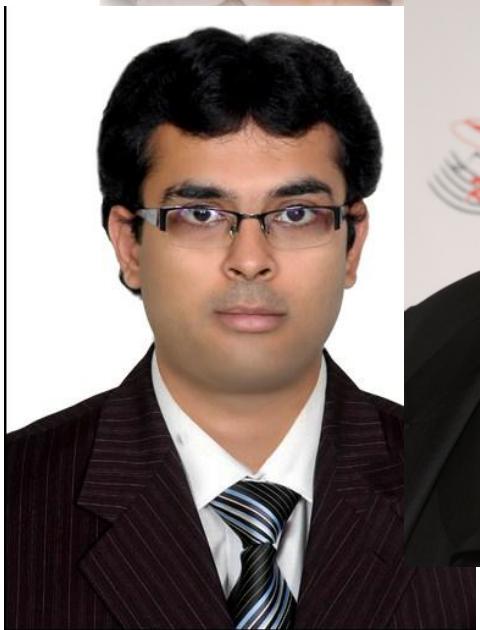


Vision: Detect Digits

0 0 0 0 0 0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2 2 2 2 2 3
3 3 3 3 3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 8 9 9 9 9



Vision: Detect Faces



Piech, CS106A, Stanford University



Vision is Hard

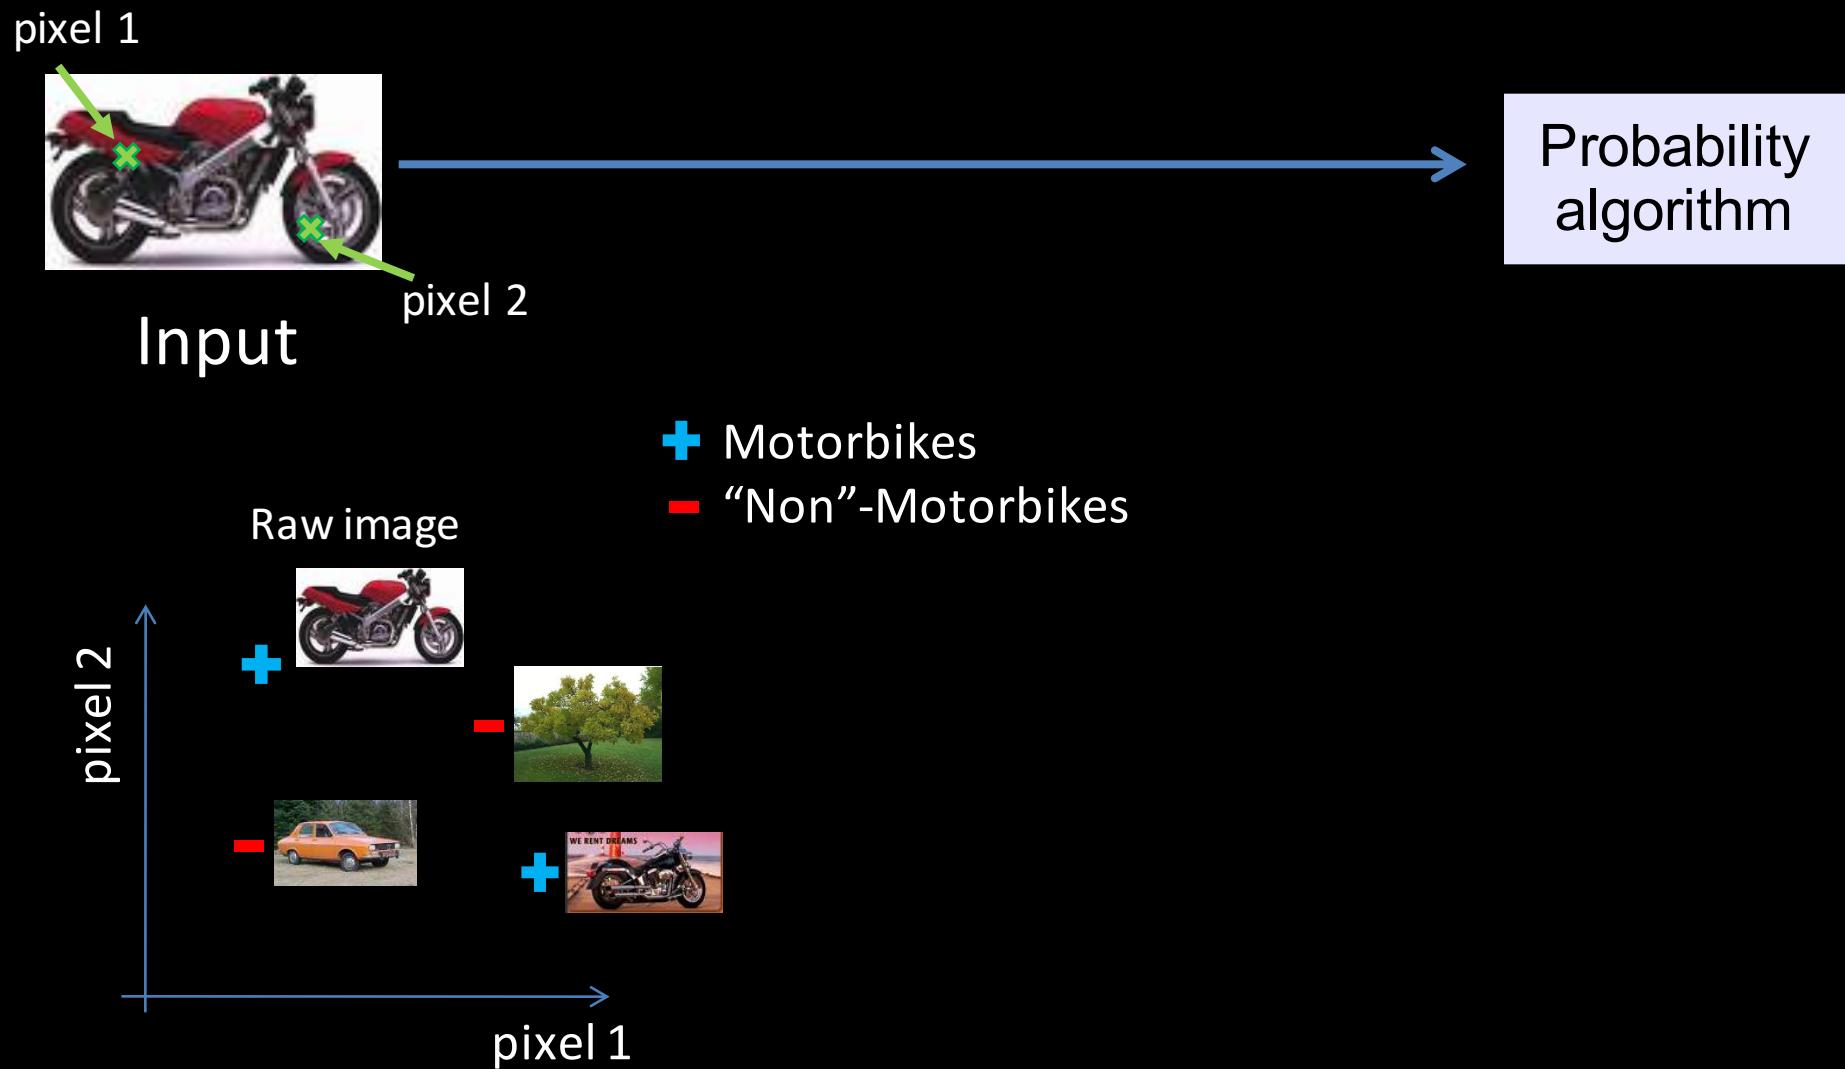
You see this:



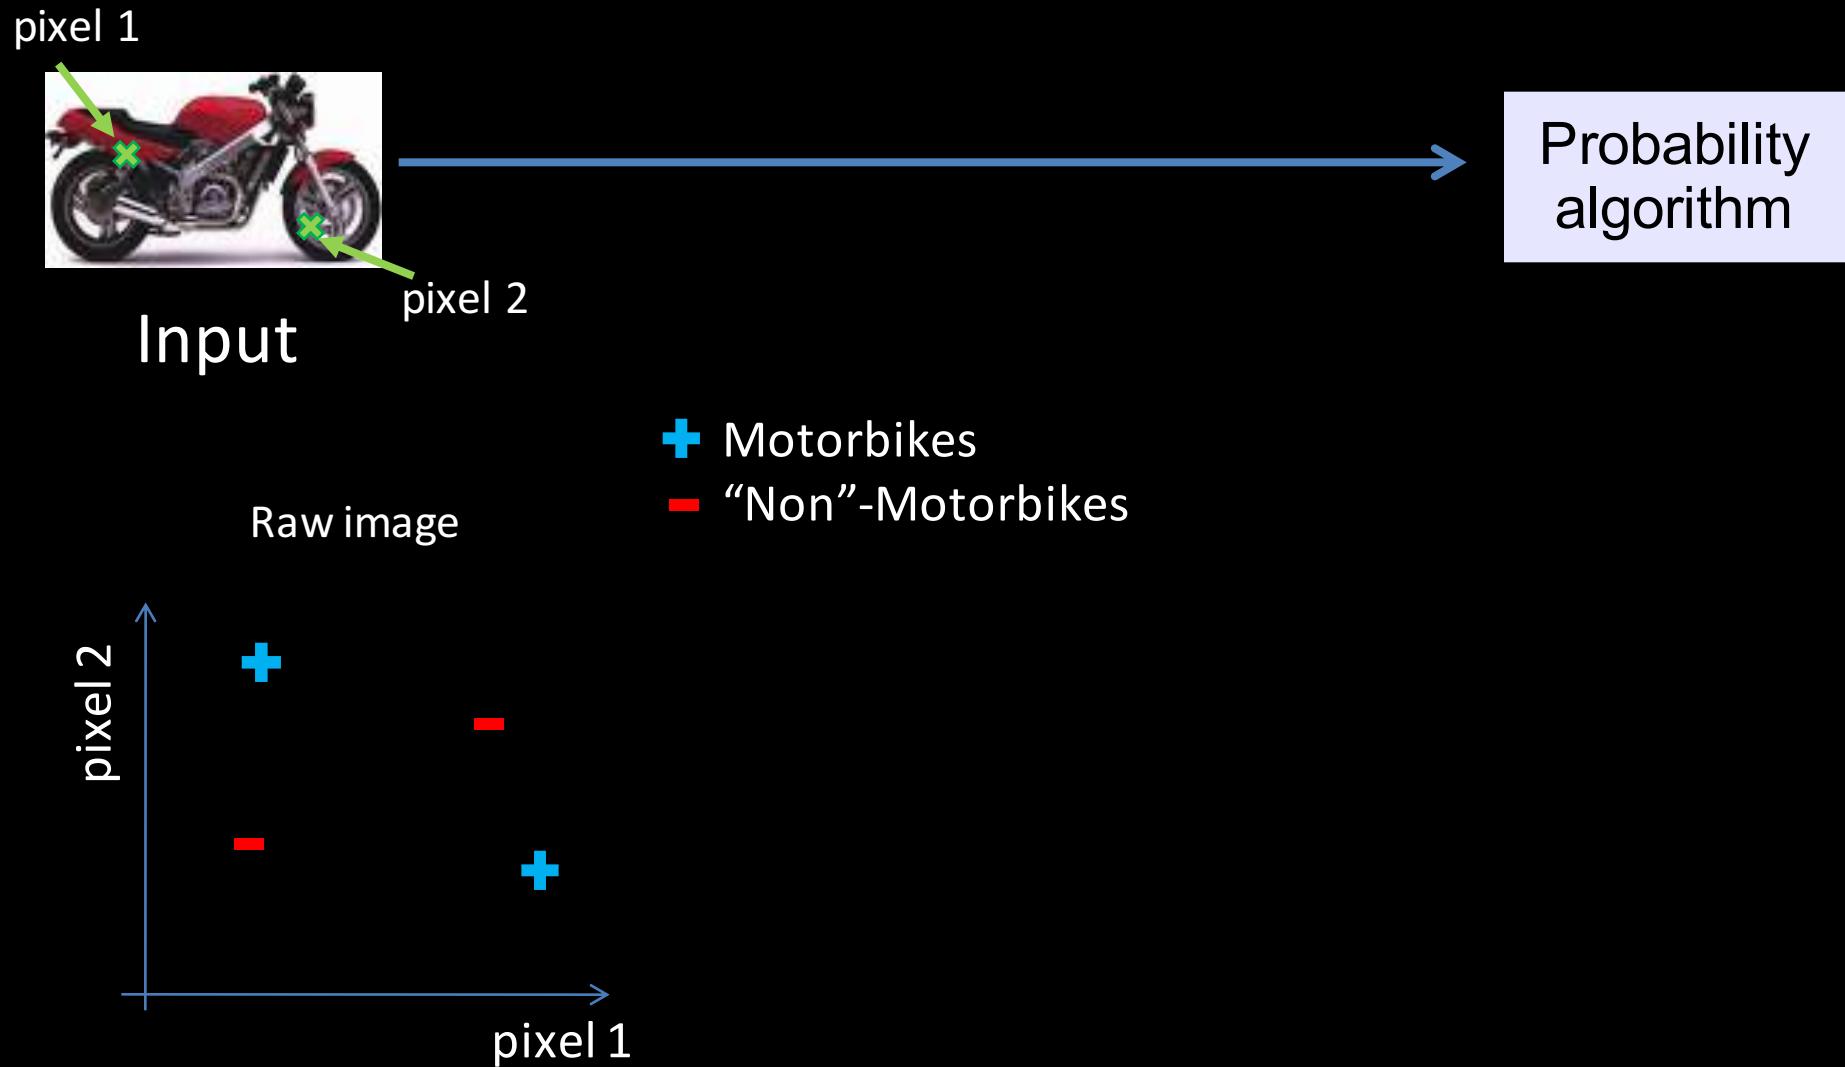
But the camera sees this:

194	210	201	212	199	213	215	195	178	158	182	209
180	189	190	221	209	205	191	167	147	115	129	163
114	126	140	188	176	165	152	140	170	106	78	88
87	103	115	154	143	142	149	153	173	101	57	57
102	112	106	131	122	138	152	147	128	84	58	66
94	95	79	104	105	124	129	113	107	87	69	67
68	71	69	98	89	92	98	95	89	88	76	67
41	56	68	99	63	45	60	82	58	76	75	65
20	43	69	75	56	41	51	73	55	70	63	44
50	50	57	69	75	75	73	74	53	68	59	37
72	59	53	66	84	92	84	74	57	72	63	42
67	61	58	65	75	78	76	73	59	75	69	50

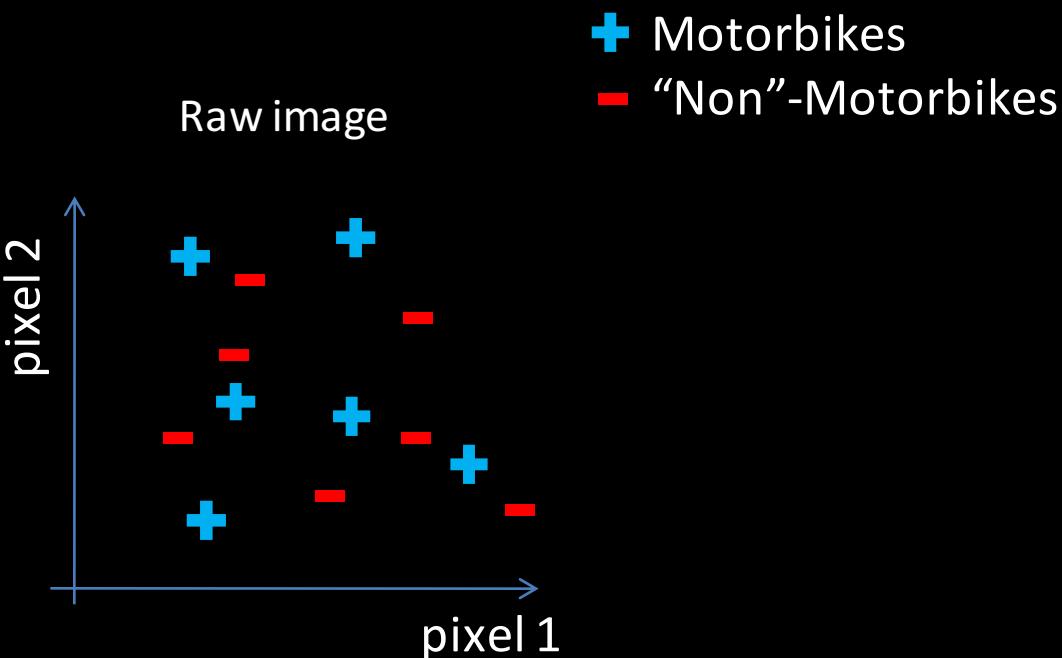
Machine Learning



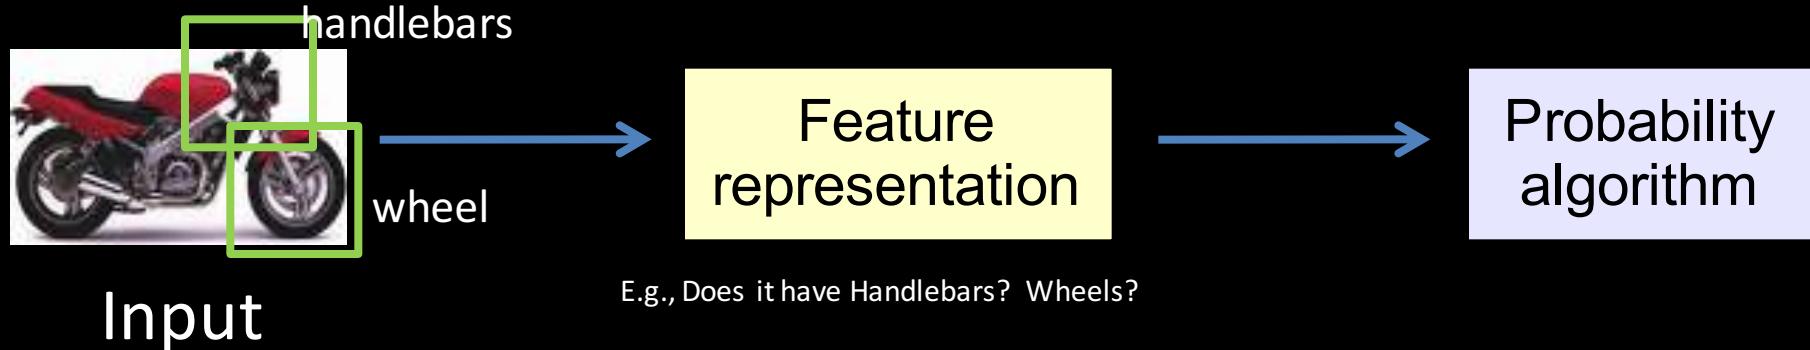
Machine Learning



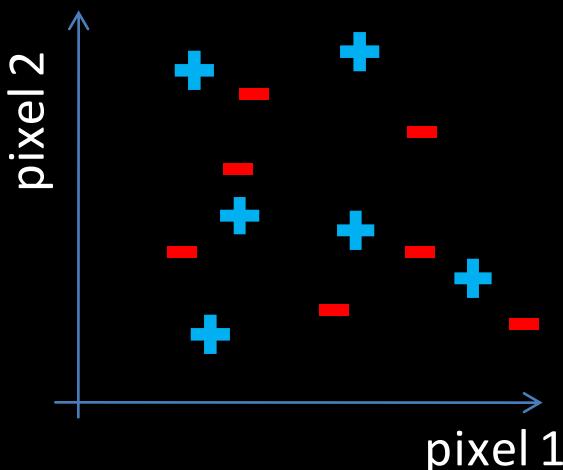
Machine Learning



Machine Learning



Raw image



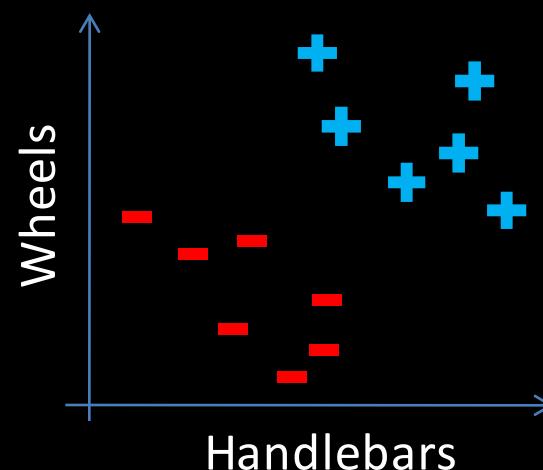
+

Motorbikes

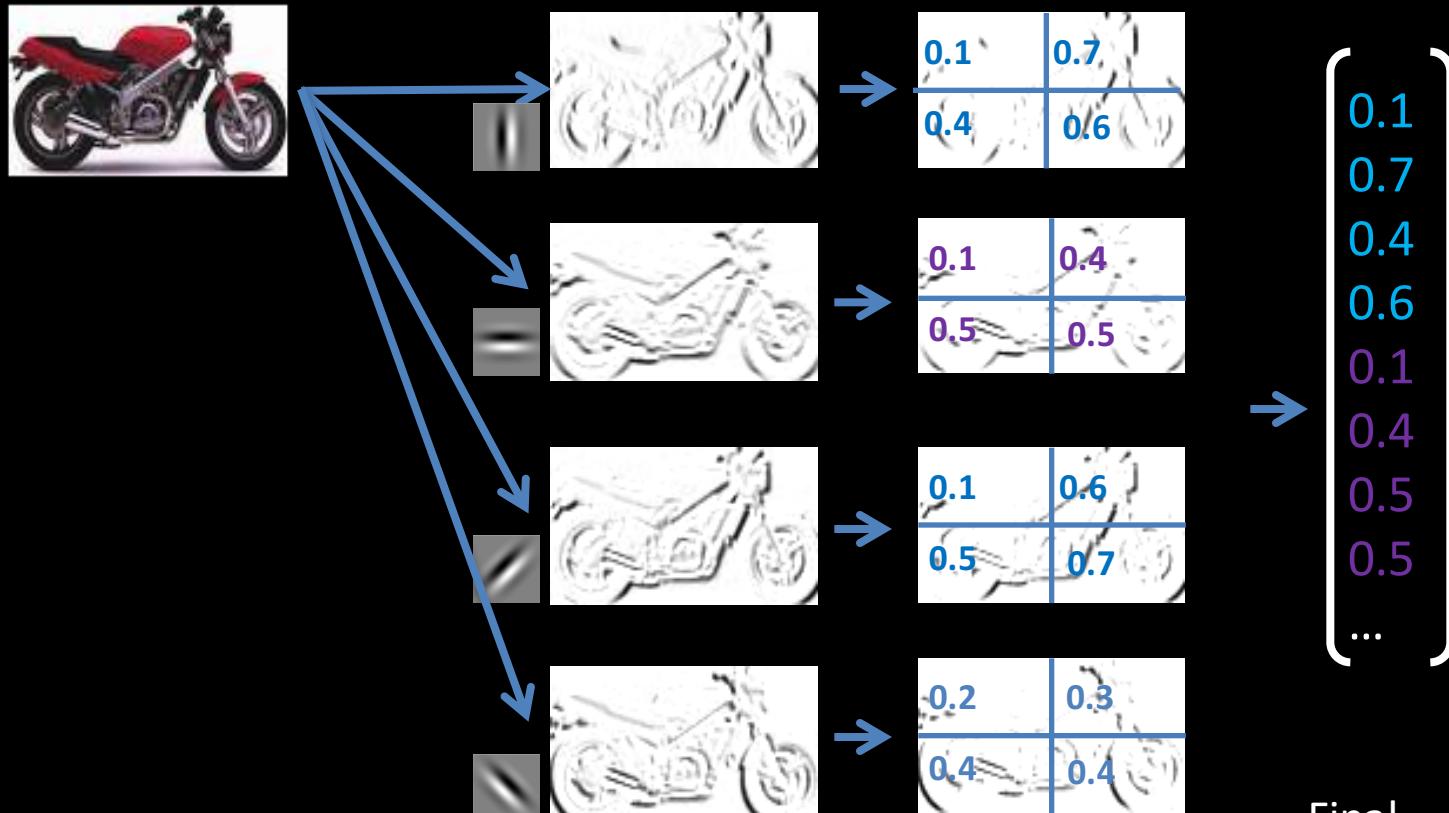
-

"Non"-Motorbikes

Features



Human Created Features



Find edges
at four
orientations

Sum up edge
strength in
each quadrant

Final
feature
vector

[Andrew Ng]

Some Great Thinkers



Daphne Koller



Mehran Sahami

Straight ML Not Perfect...

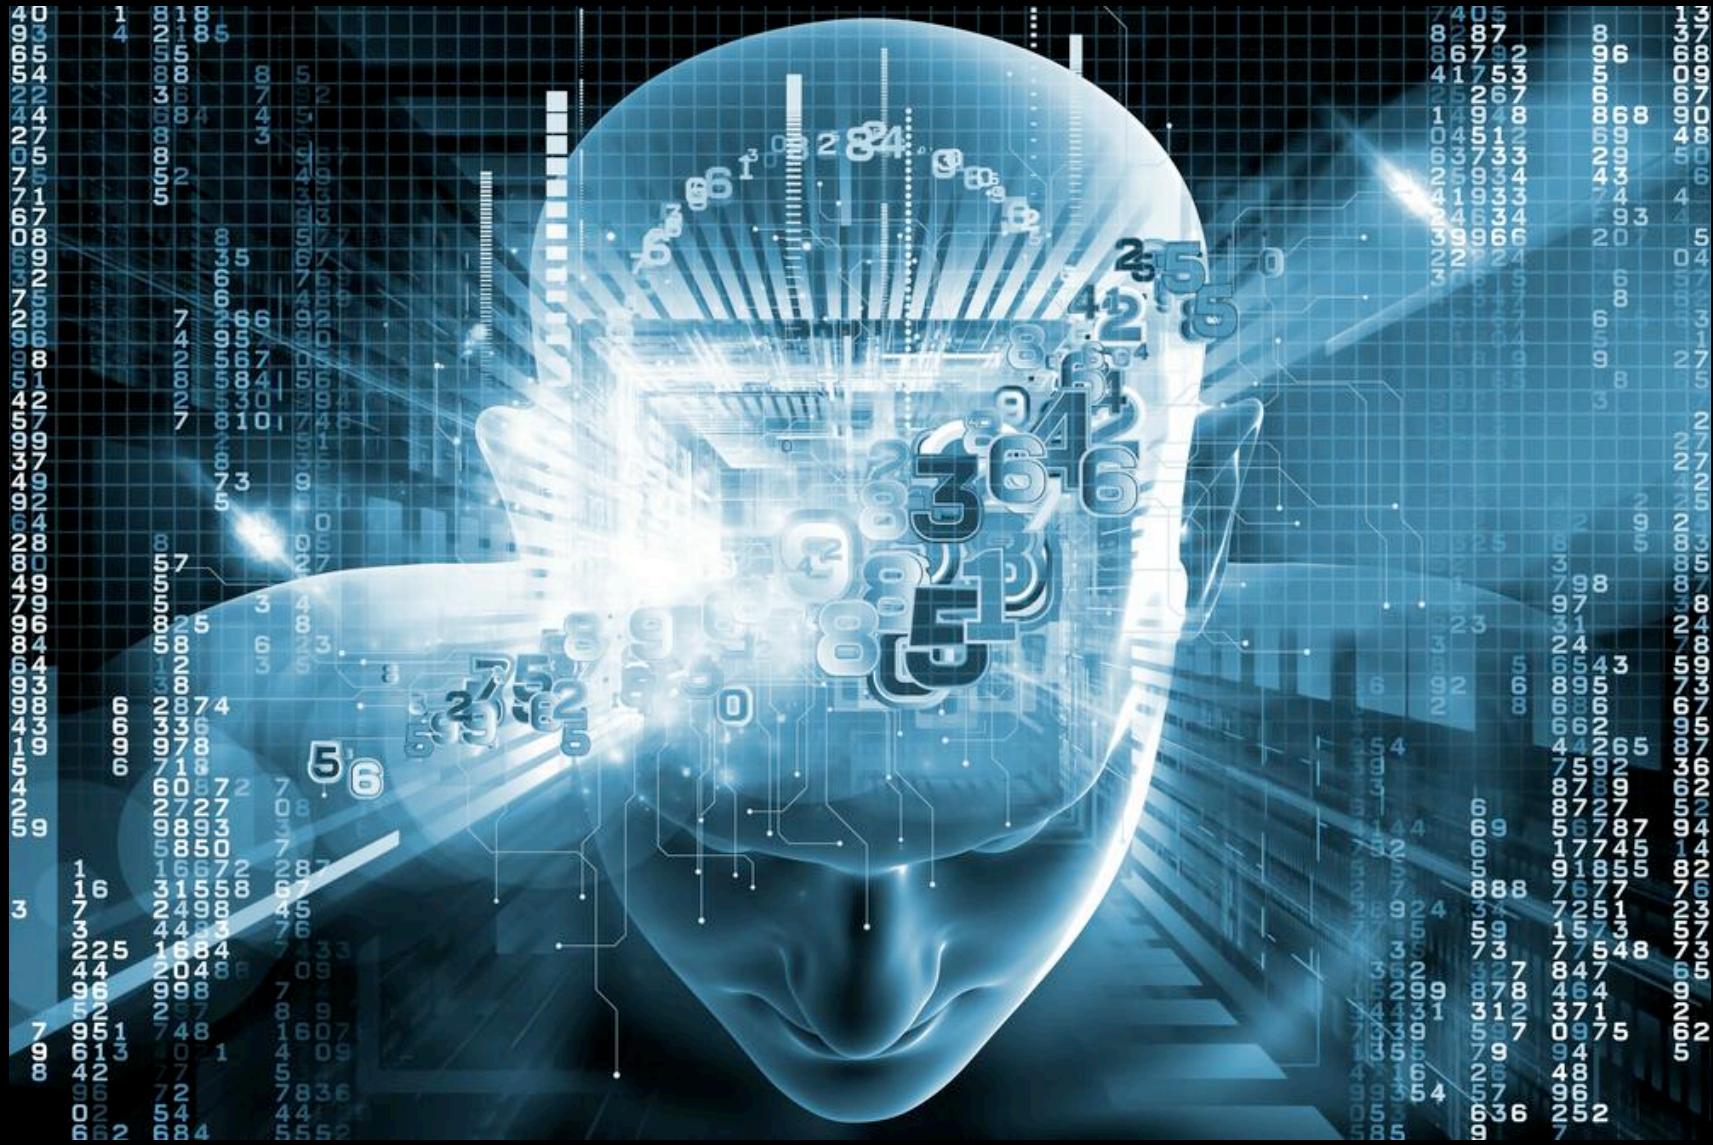


Two Great Ideas

1. Probability from Examples

2. Artificial Neurons

2. Artificial Neurons



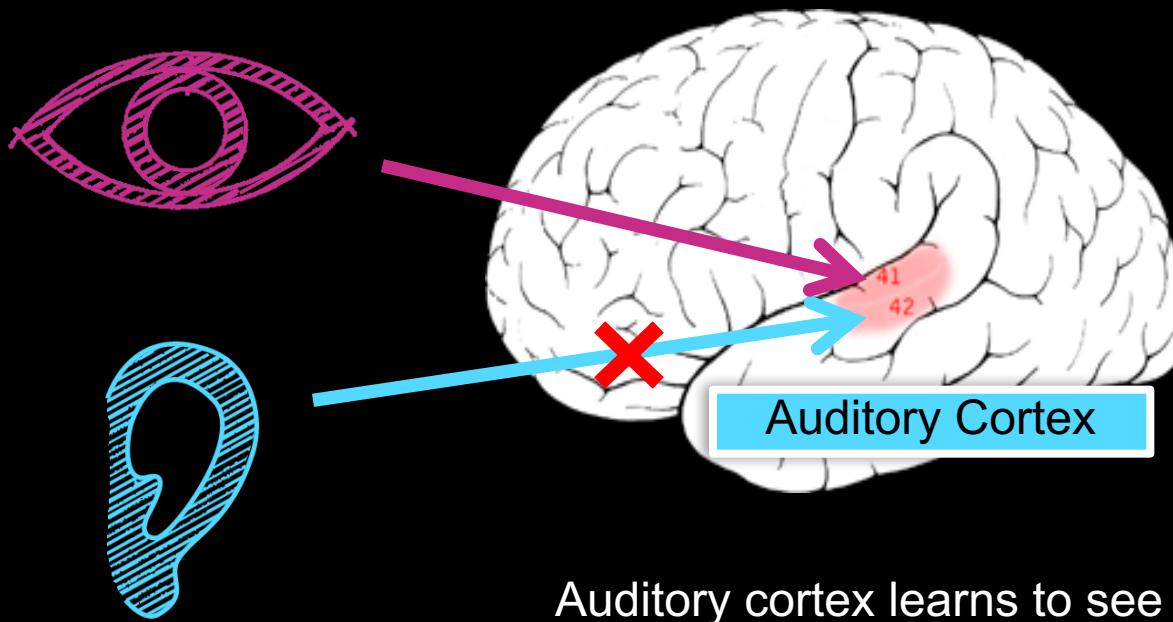
One Algorithm Hypothesis

Much of perception in the brain can be explained with a single learning algorithm.



[Andrew Ng]

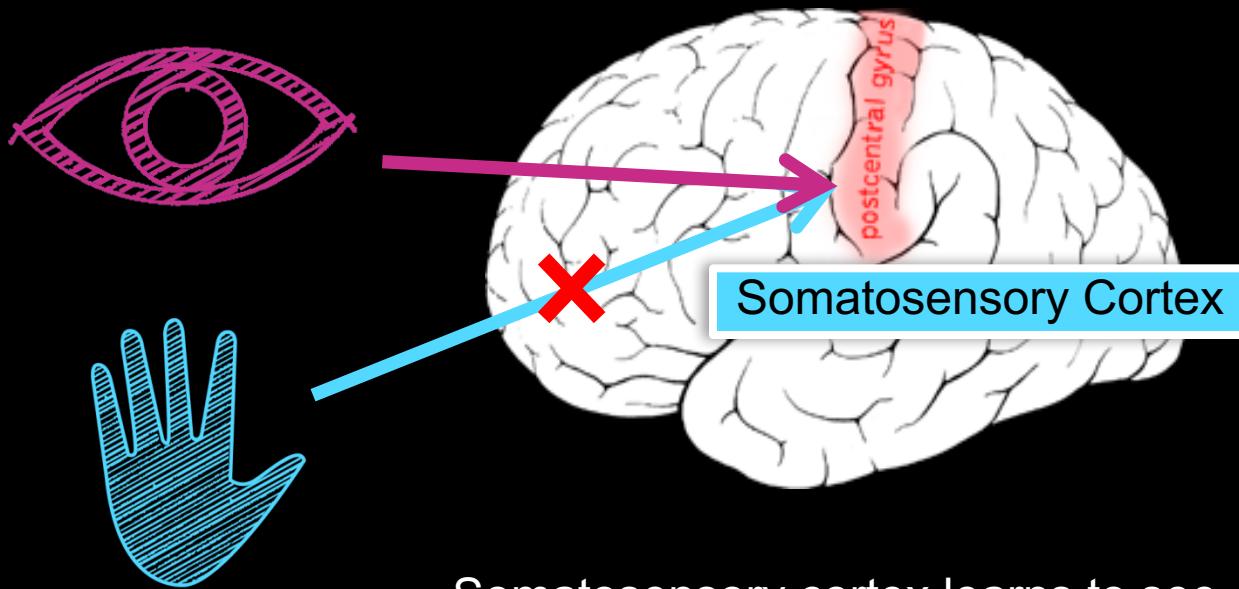
One Algorithm Hypothesis



[Roe et al., 1992]

[Andrew Ng]

One Algorithm Hypothesis

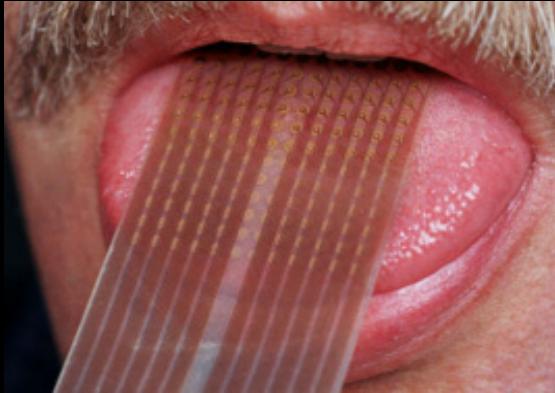


Somatosensory cortex learns to see

[Metin & Frost, 1989]

[Andrew Ng]

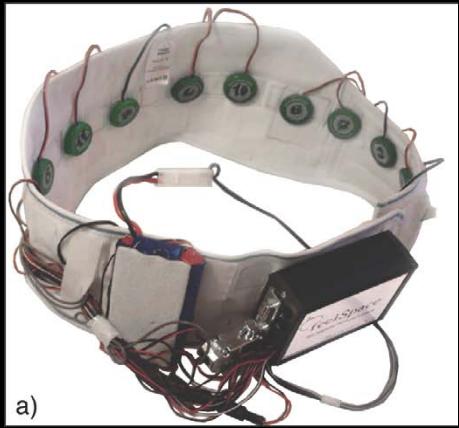
Sensor Representations



Seeing with your tongue



Human echolocation (sonar)

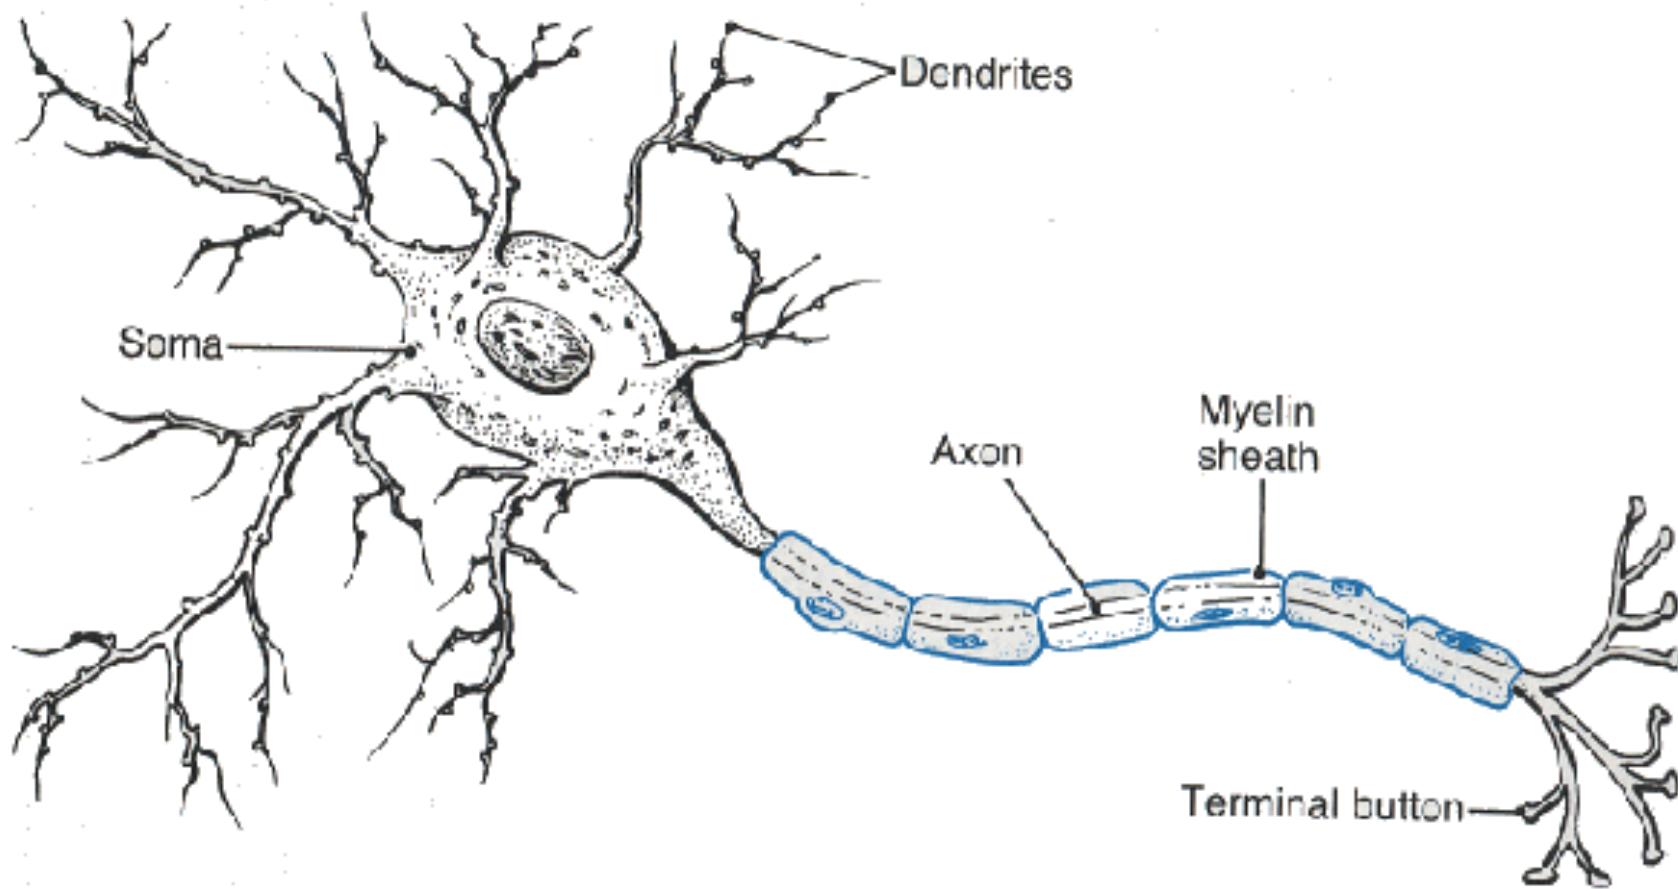


Haptic belt: Direction sense

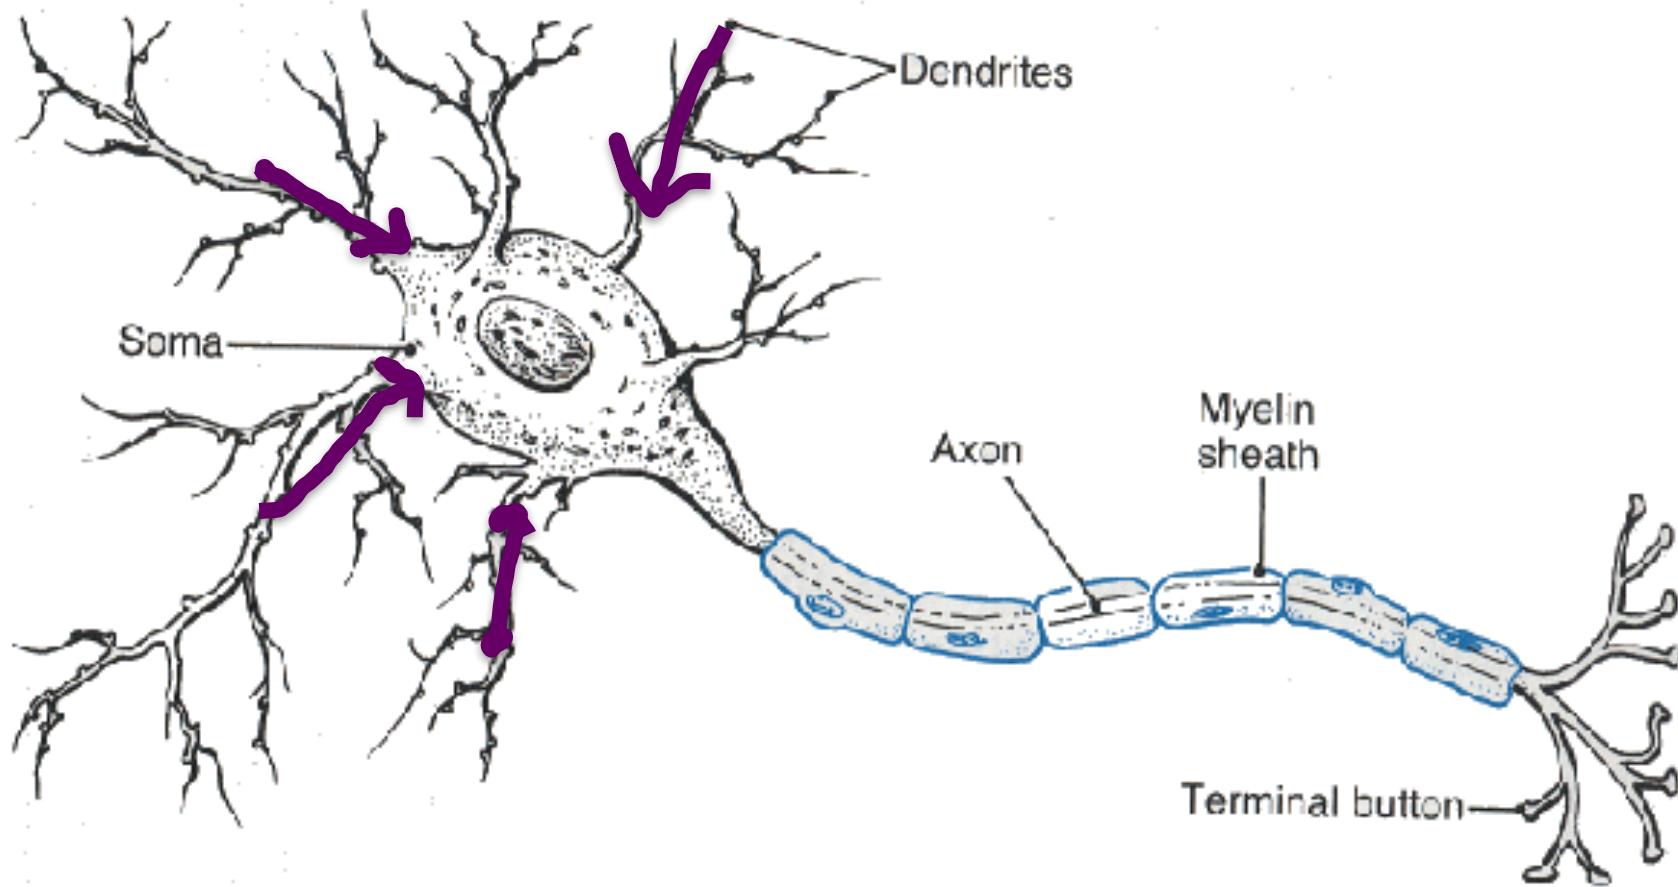


Implanting a 3rd eye

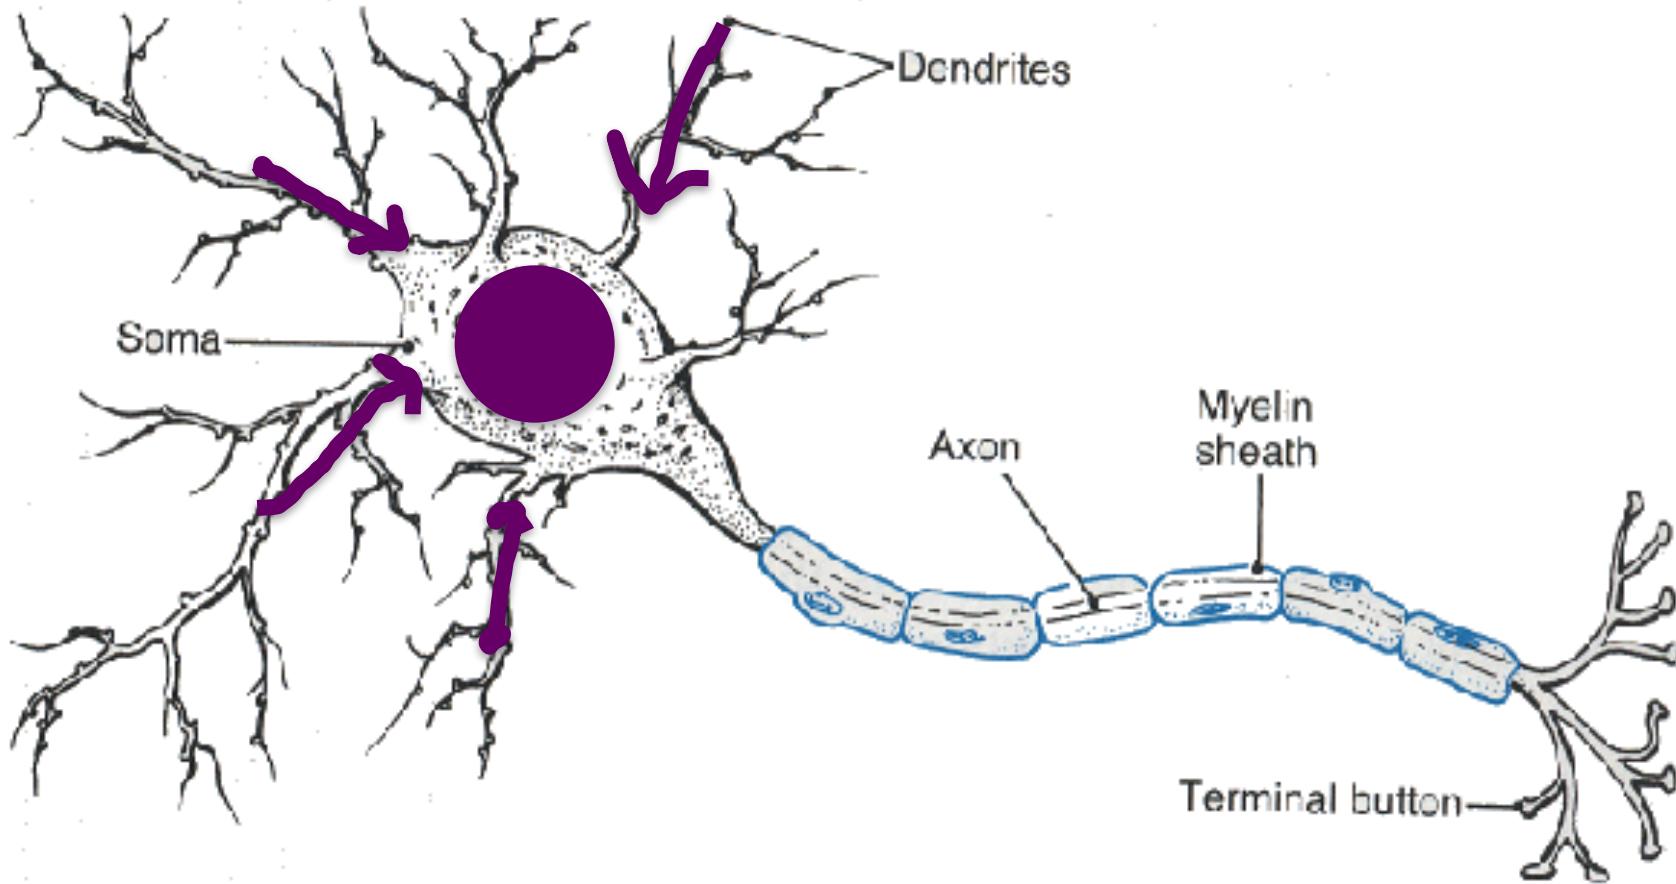
Neuron



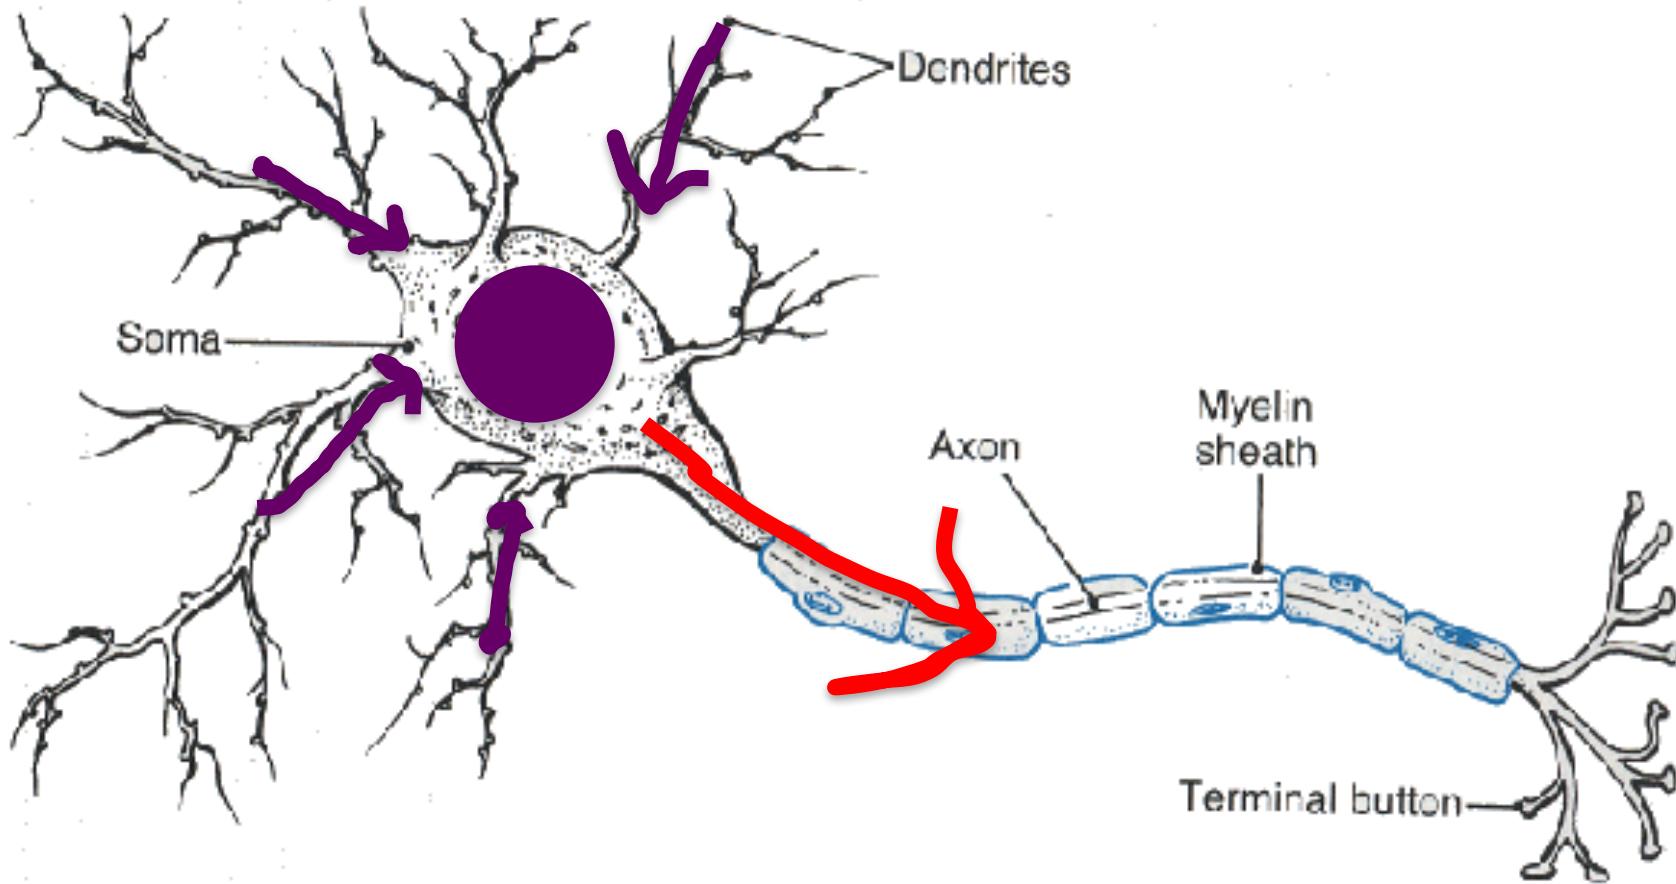
Neuron



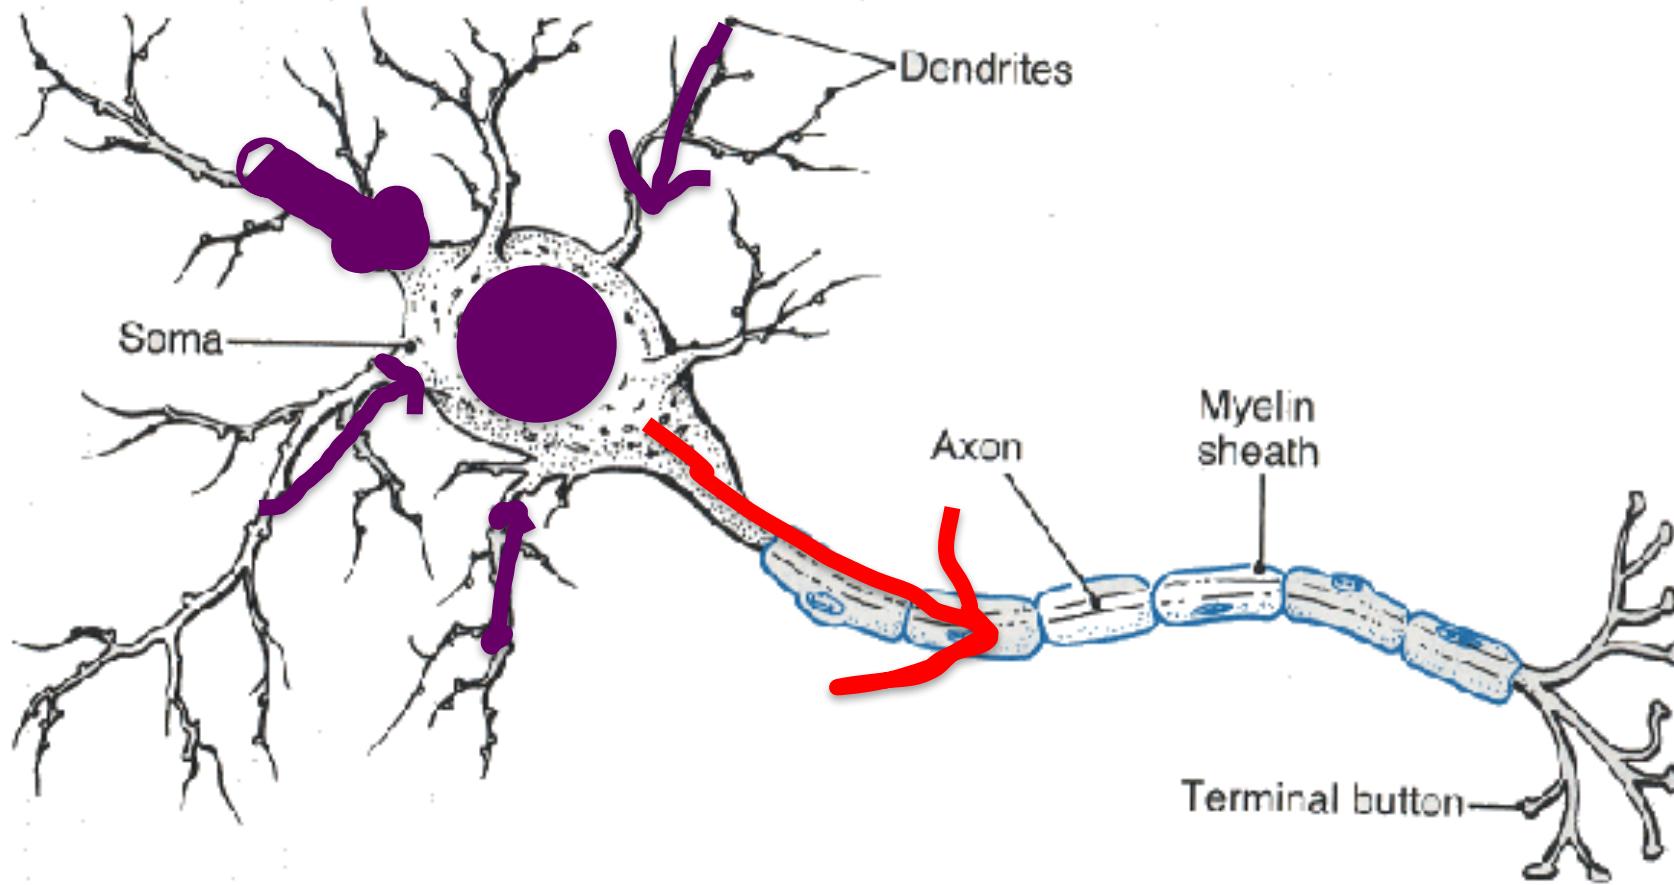
Neuron



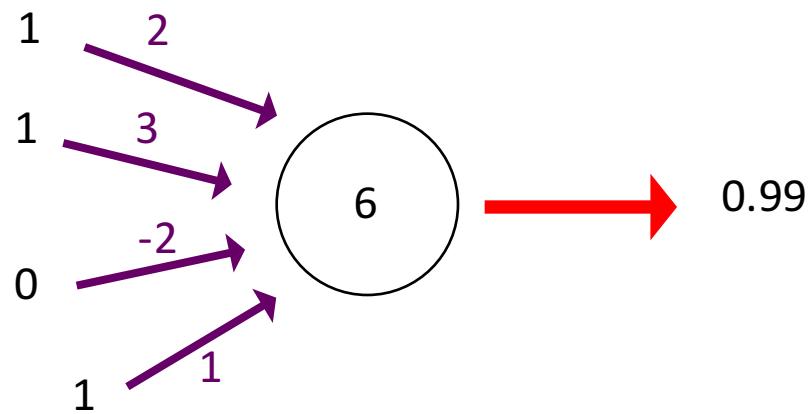
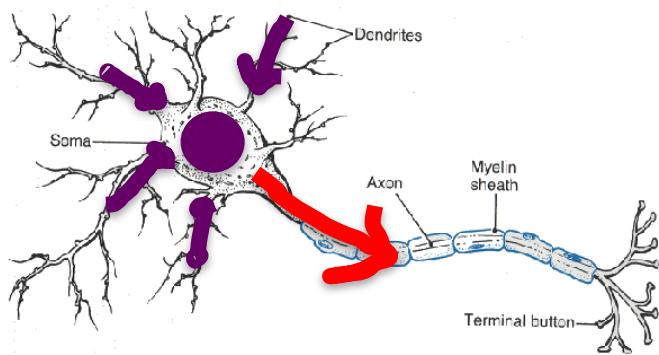
Neuron



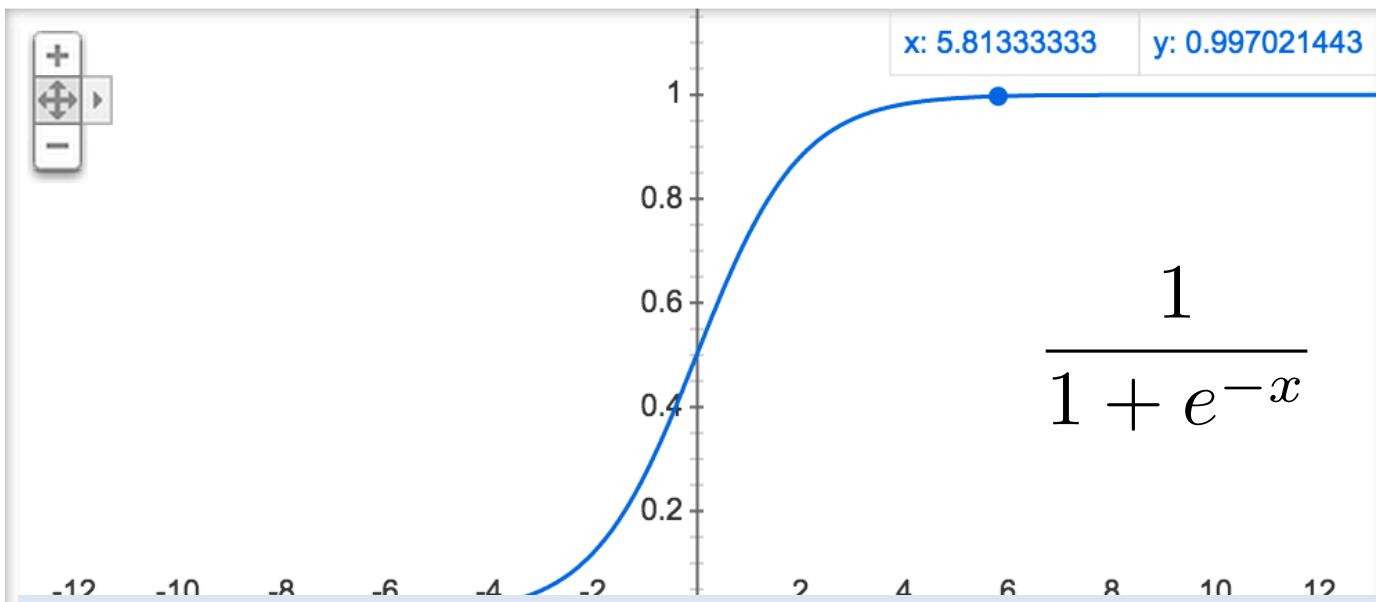
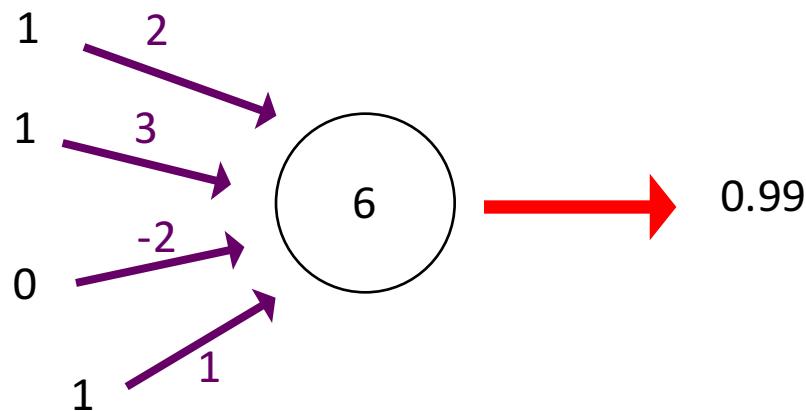
Some Inputs are More Important



Artificial Neuron



Sigmoid Function



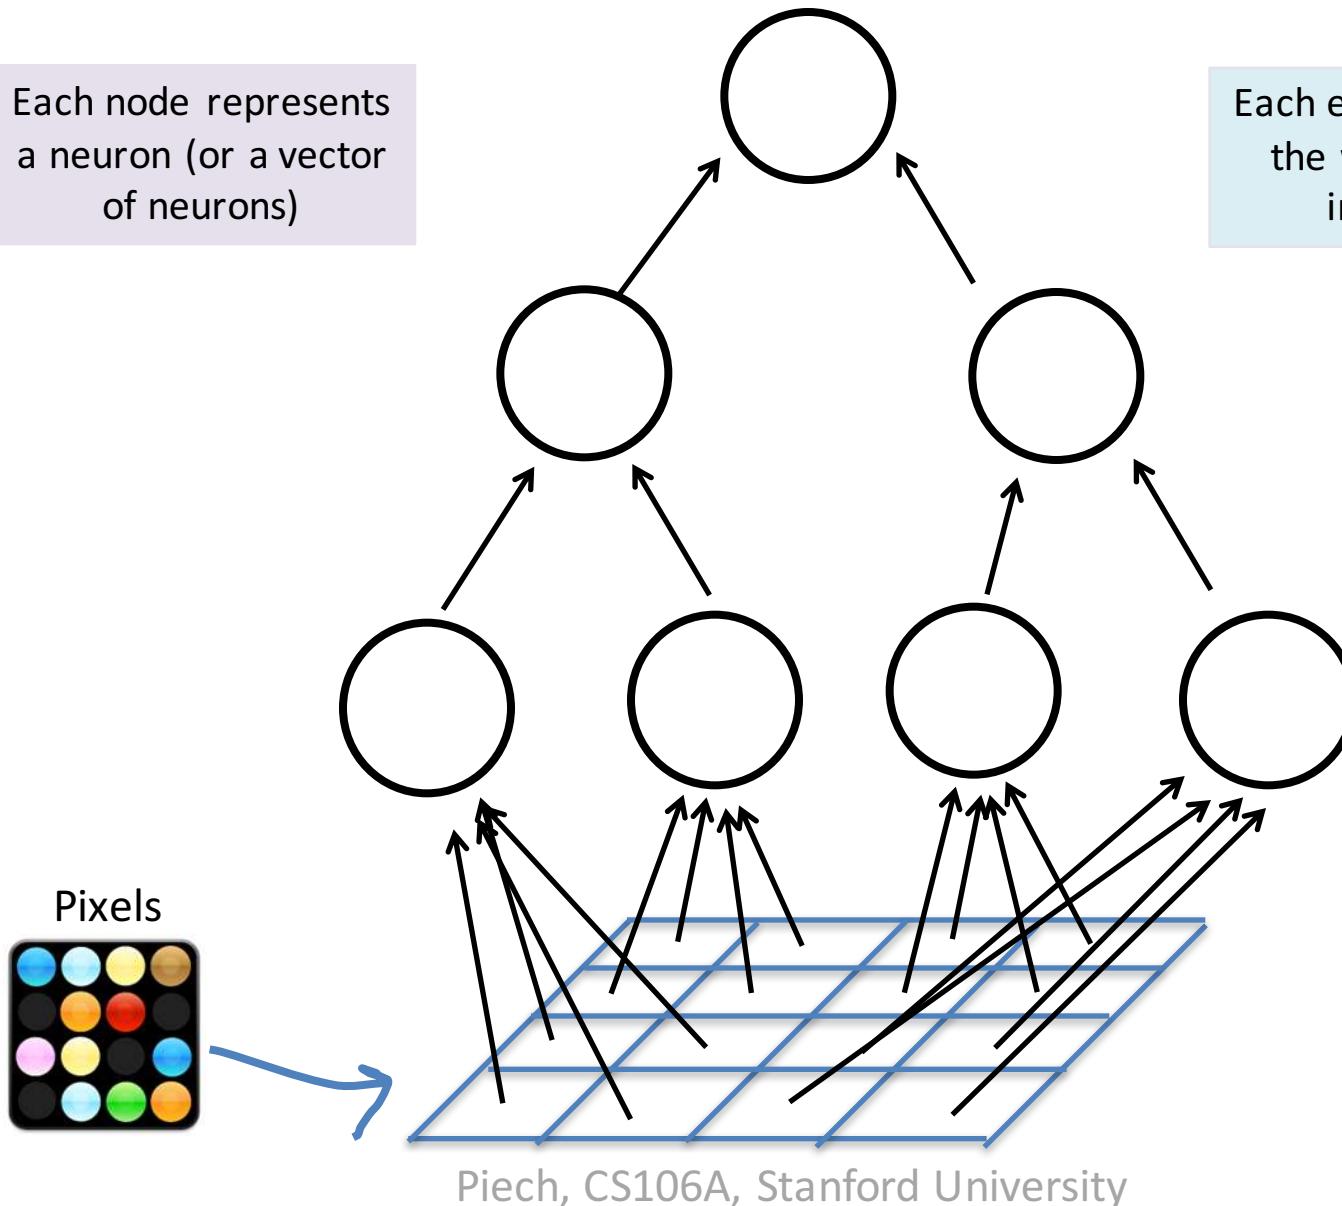
An artificial neuron is like a little probability calculator



Neural Network

Each node represents
a neuron (or a vector
of neurons)

Each edge represents
the weight of the
interaction



Piech, CS106A, Stanford University

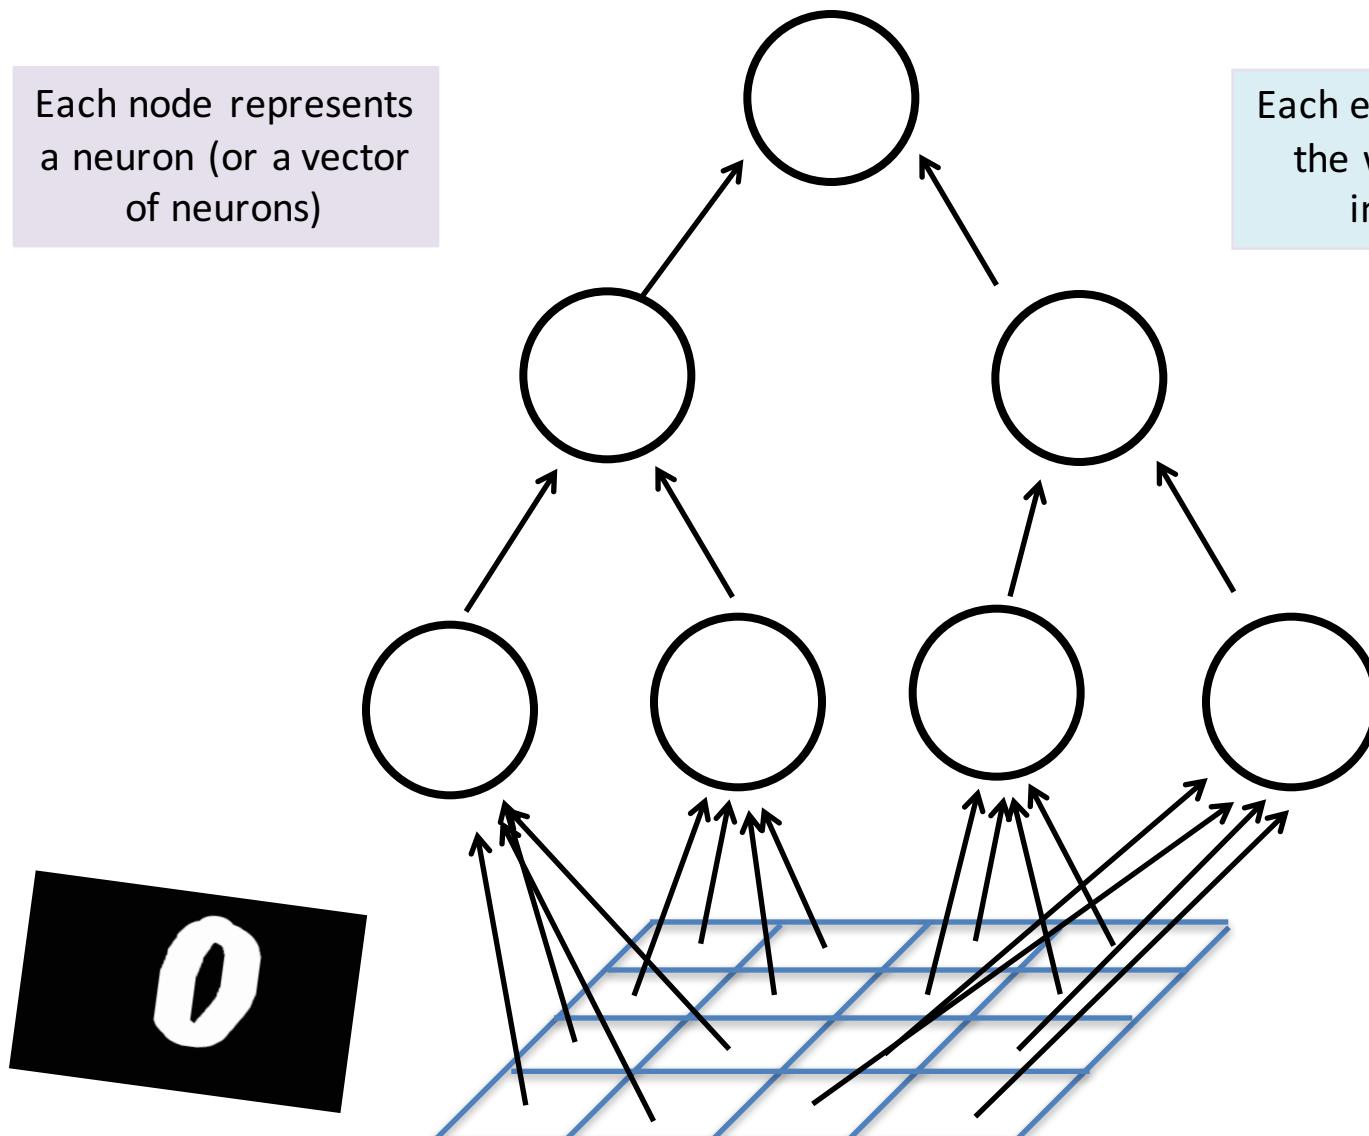


Forward Pass...

Forward Pass

Each node represents
a neuron (or a vector
of neurons)

Each edge represents
the weight of the
interaction



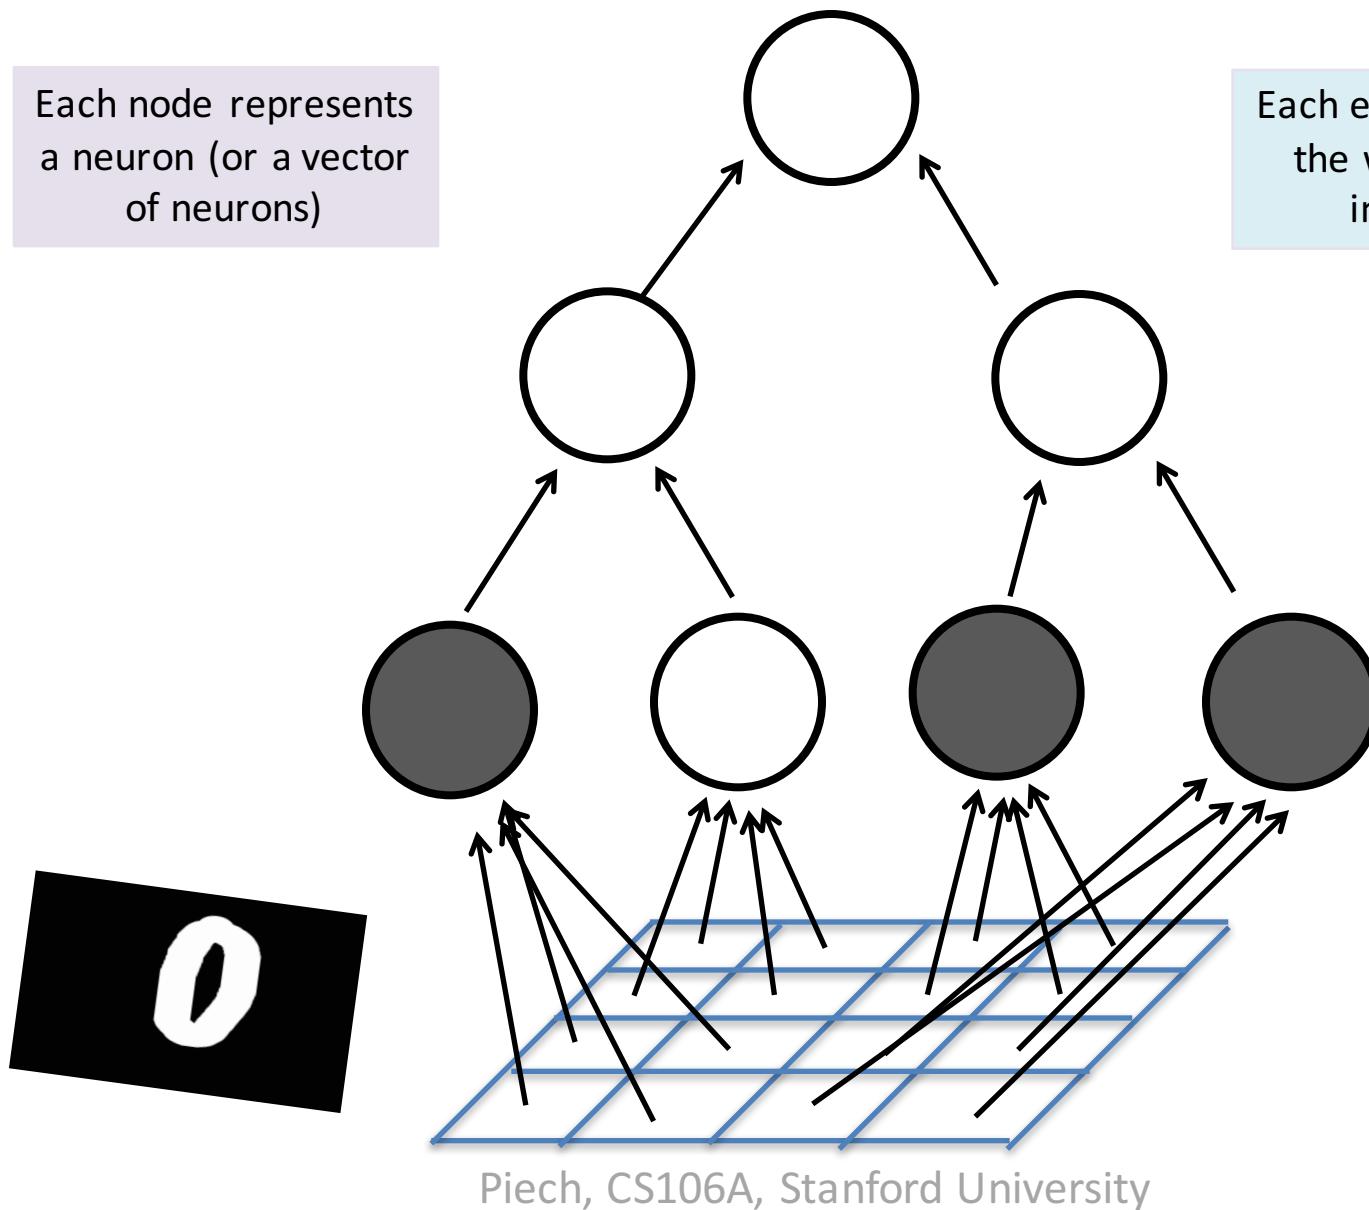
Piech, CS106A, Stanford University



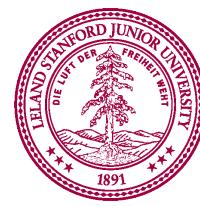
Forward Pass

Each node represents
a neuron (or a vector
of neurons)

Each edge represents
the weight of the
interaction



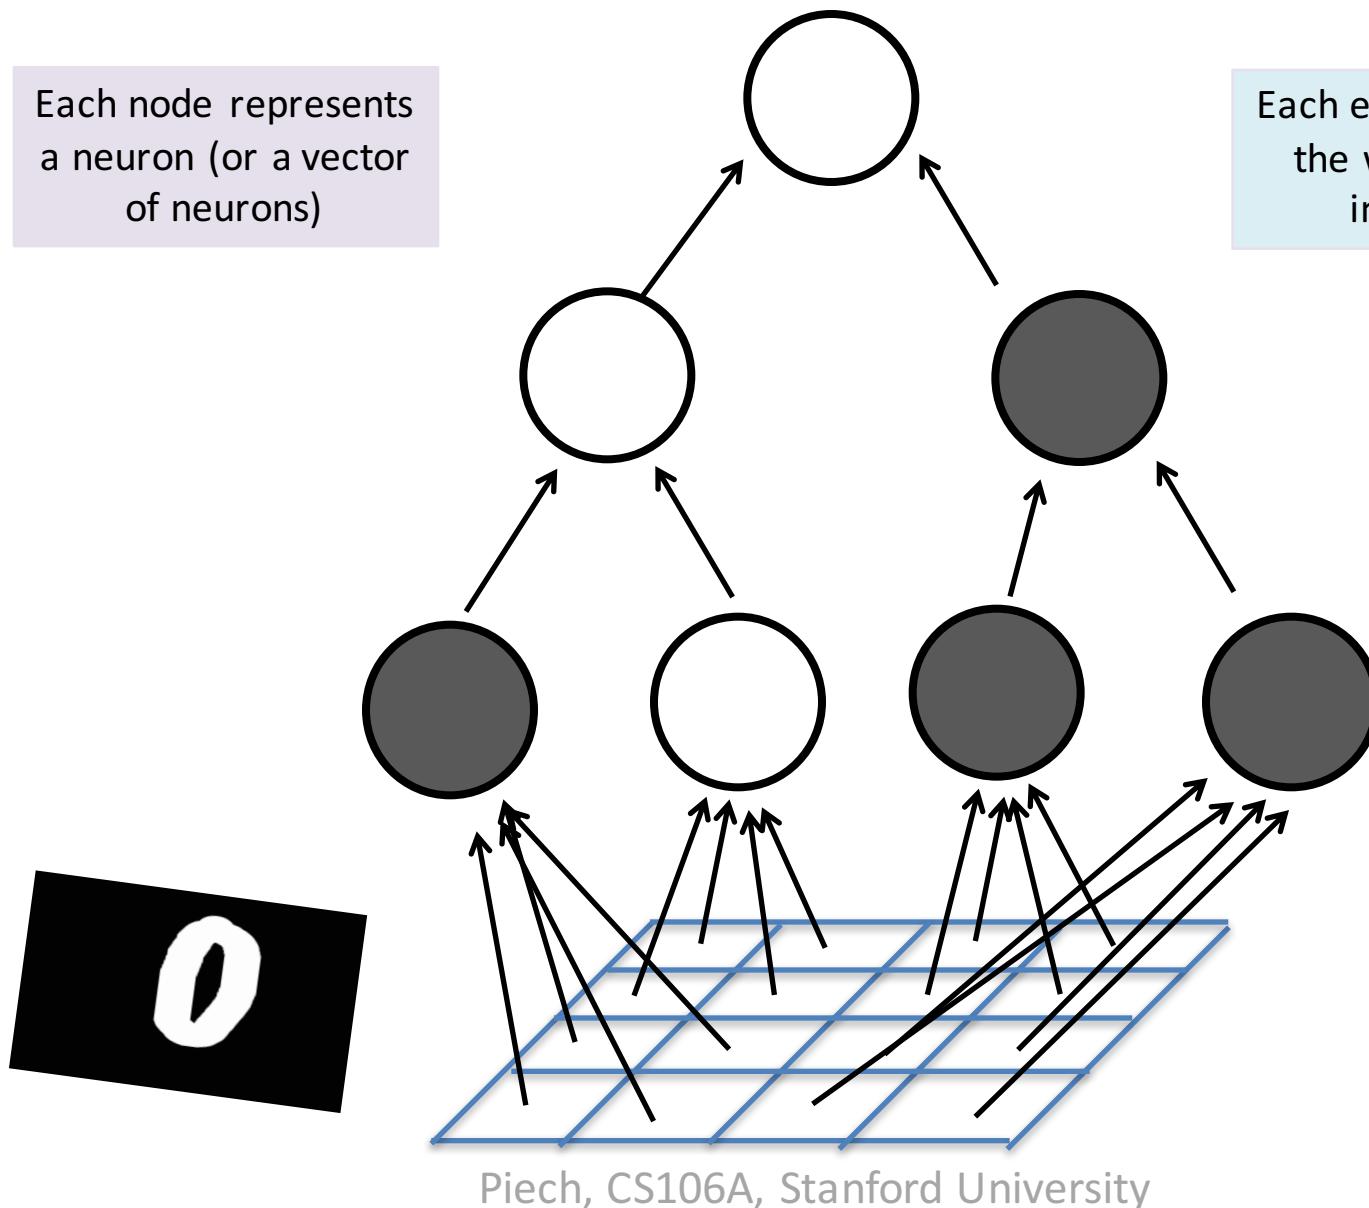
Piech, CS106A, Stanford University



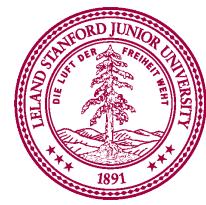
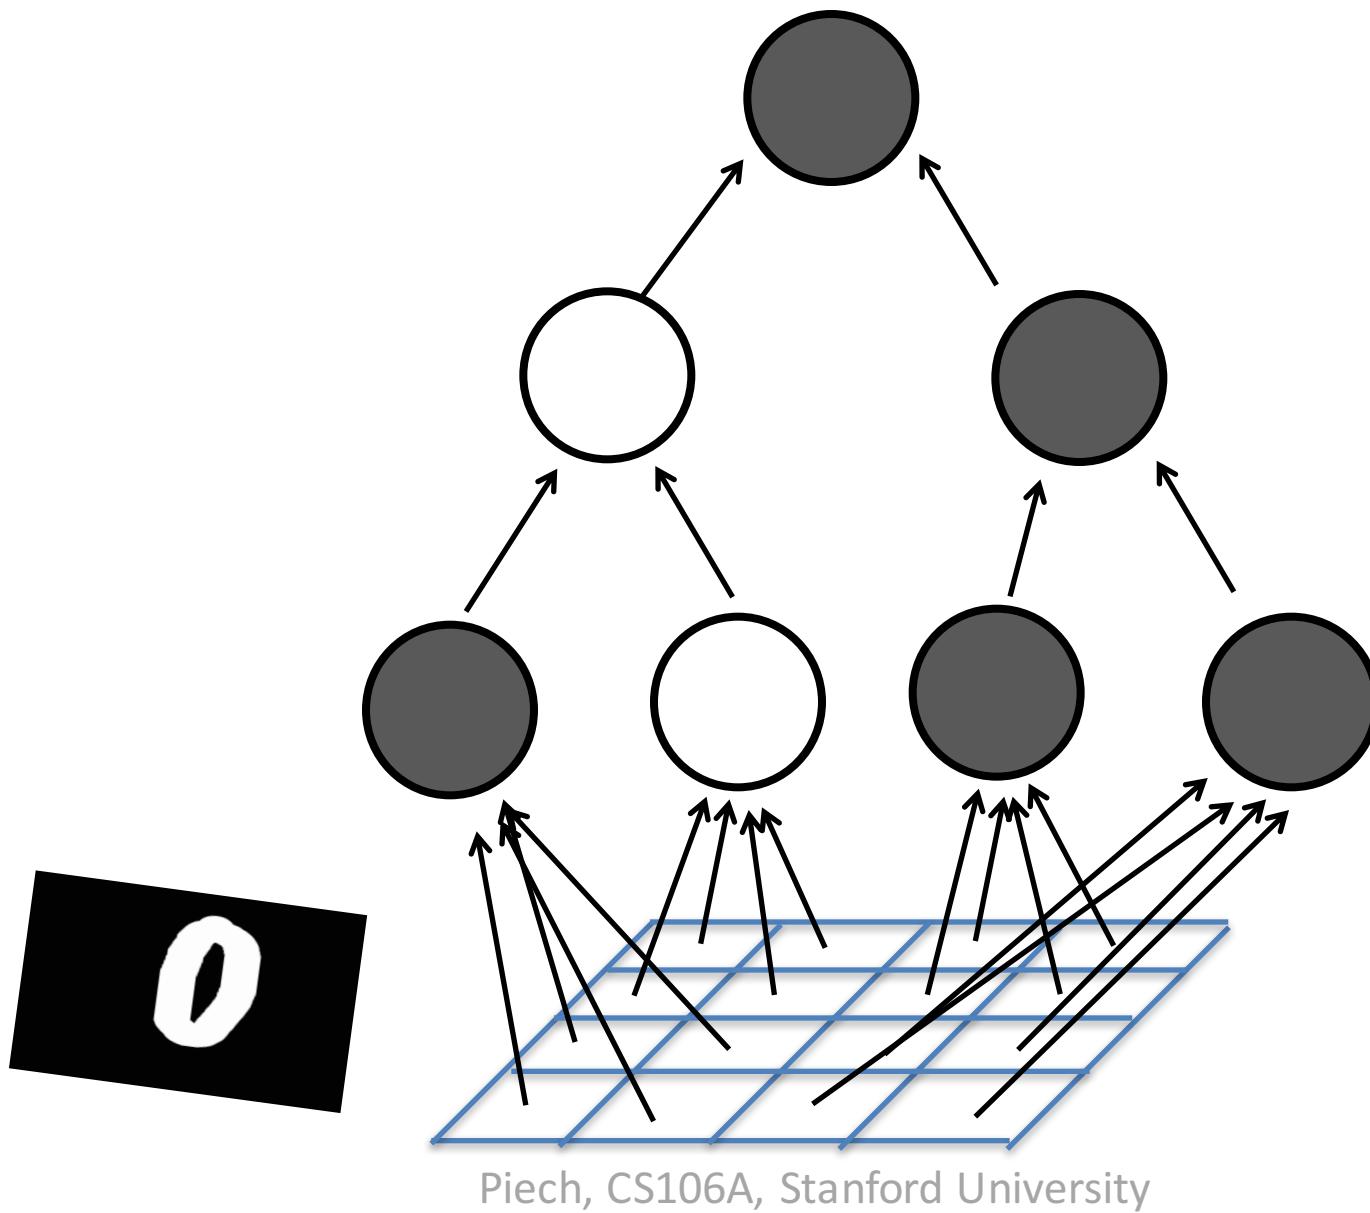
Forward Pass

Each node represents
a neuron (or a vector
of neurons)

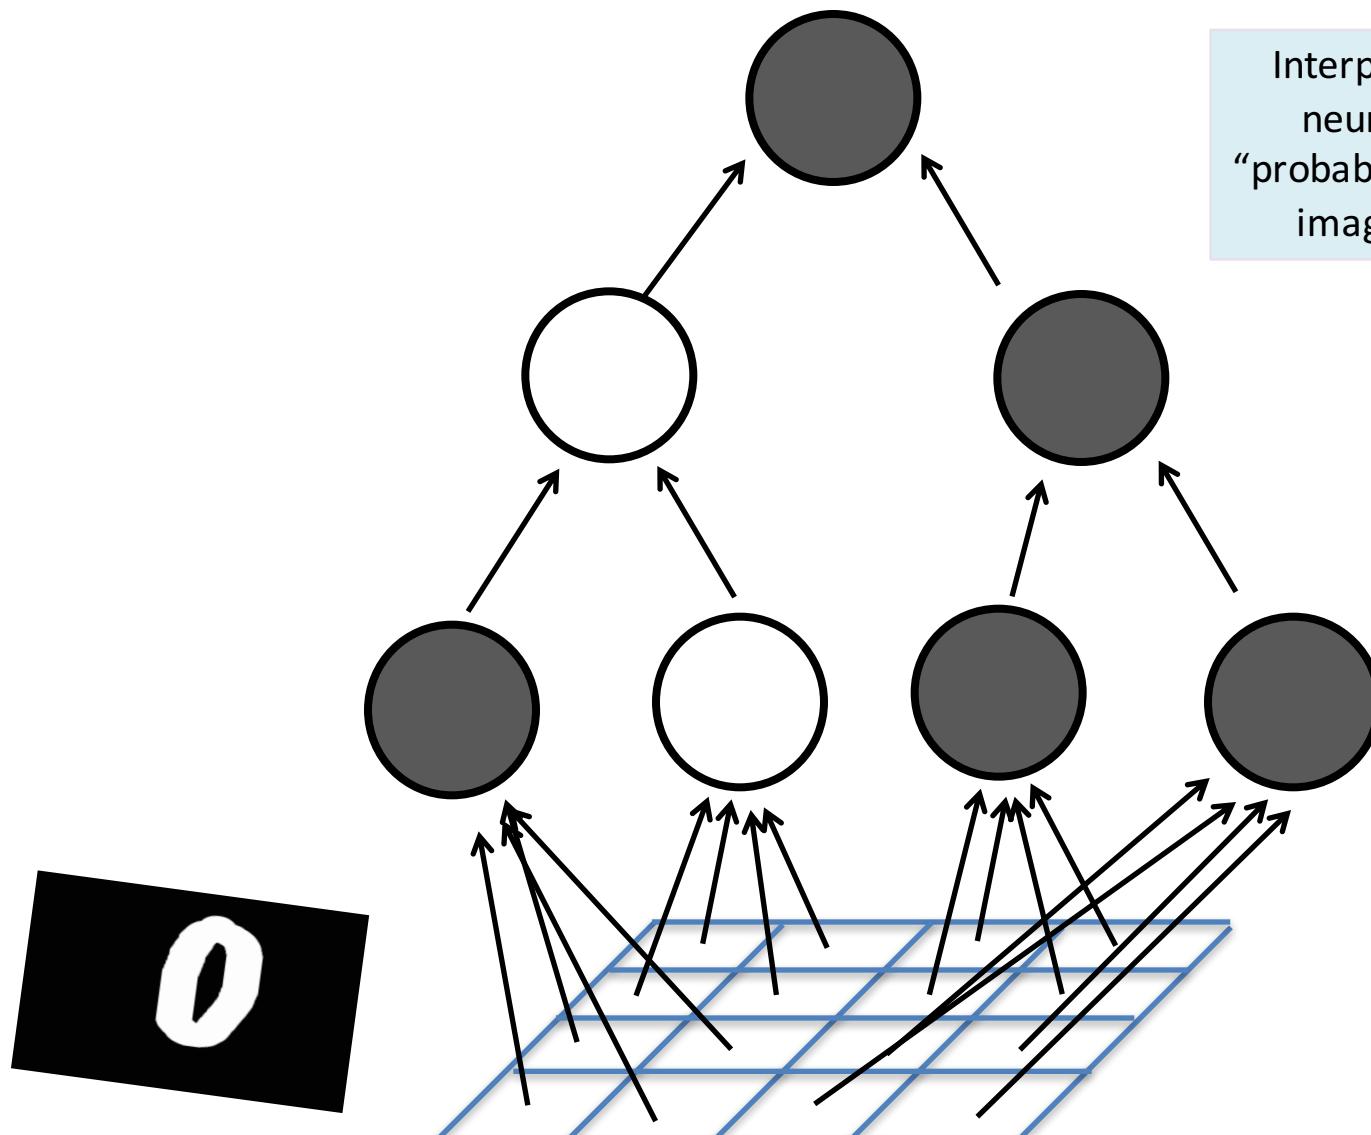
Each edge represents
the weight of the
interaction



Forward Pass



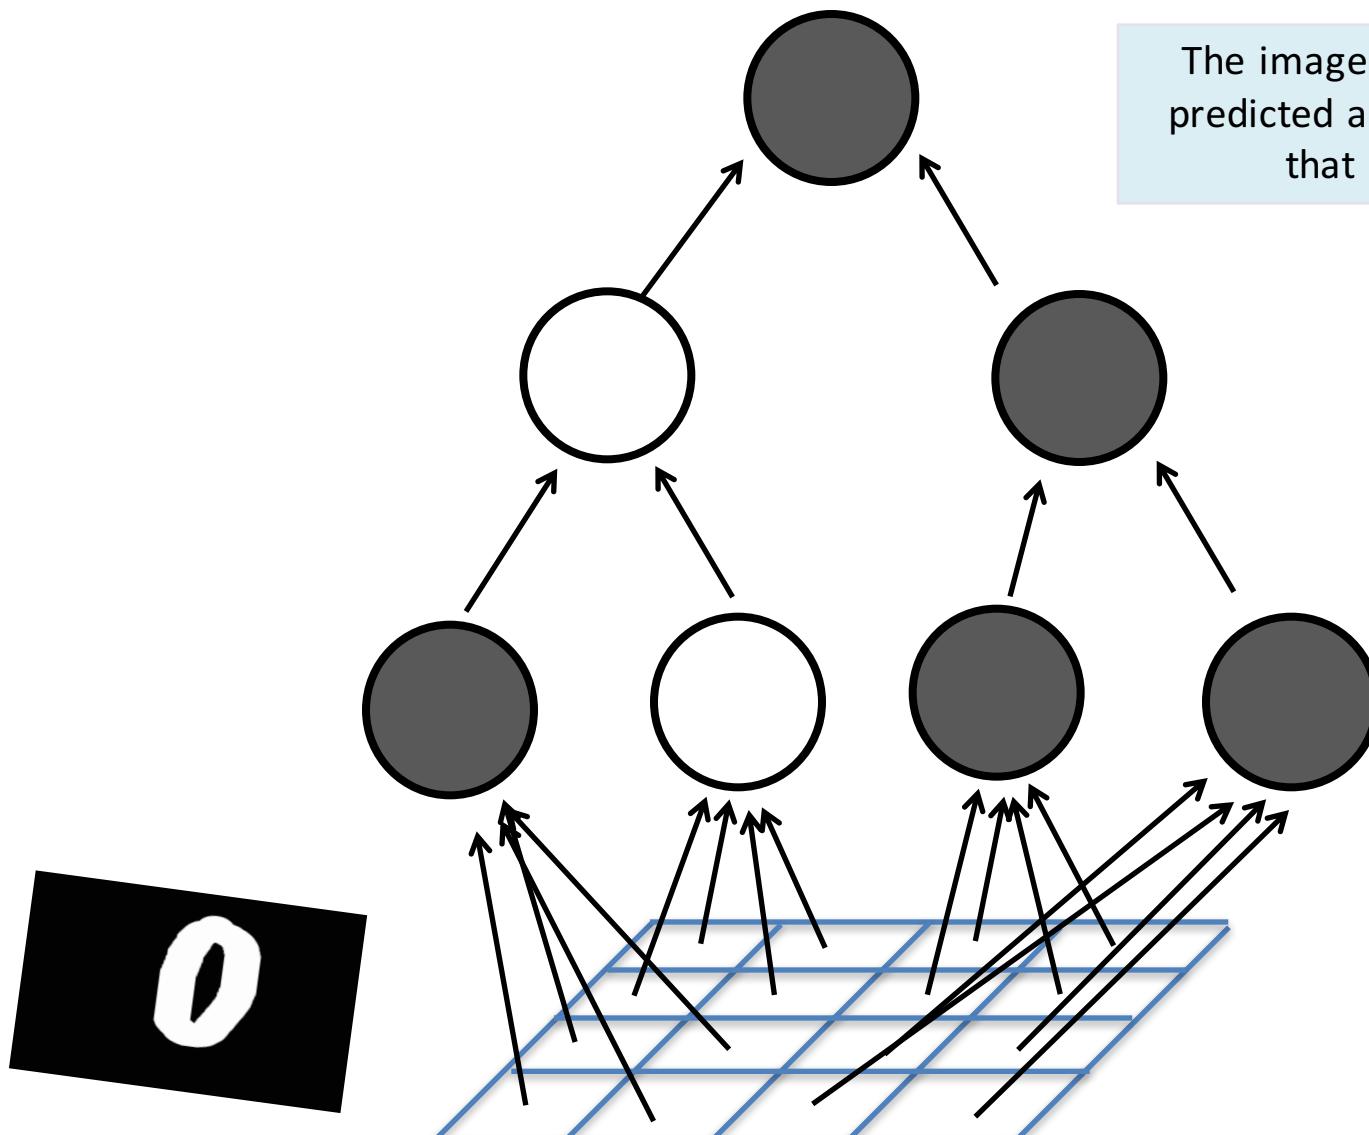
Forward Pass



Interpret the last
neuron as the
“probability” that the
image is of a 1



Backward Pass



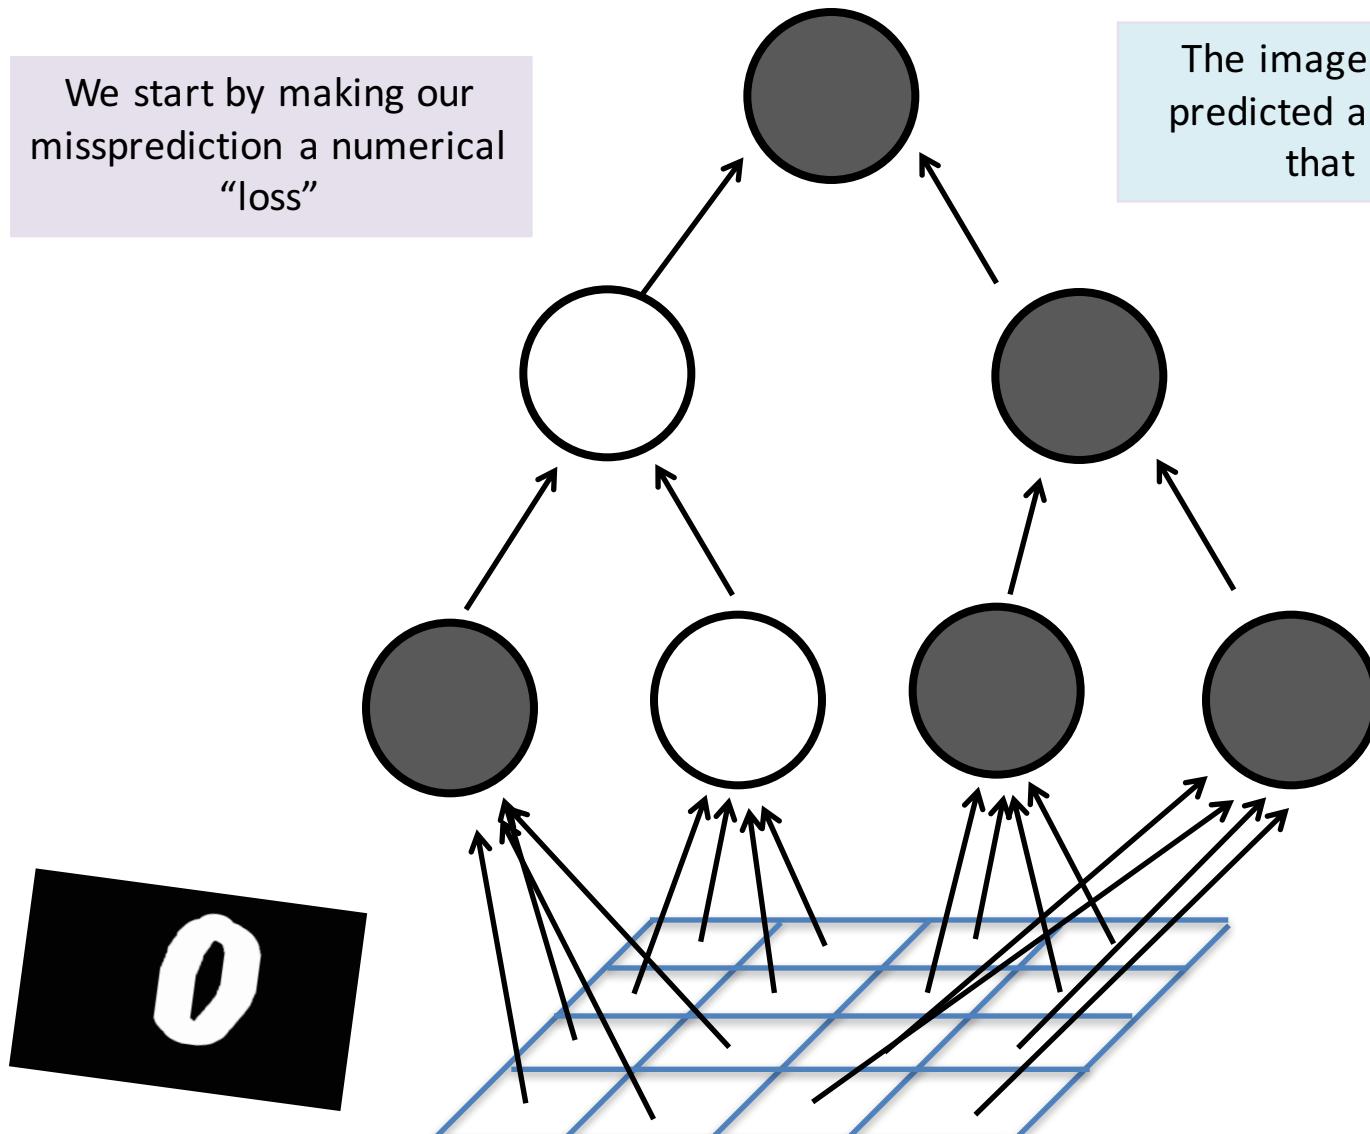
Piech, CS106A, Stanford University



Backward Pass

We start by making our missprediction a numerical “loss”

The image had a 0 but we predicted a high probability that it was a 1



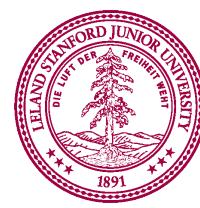
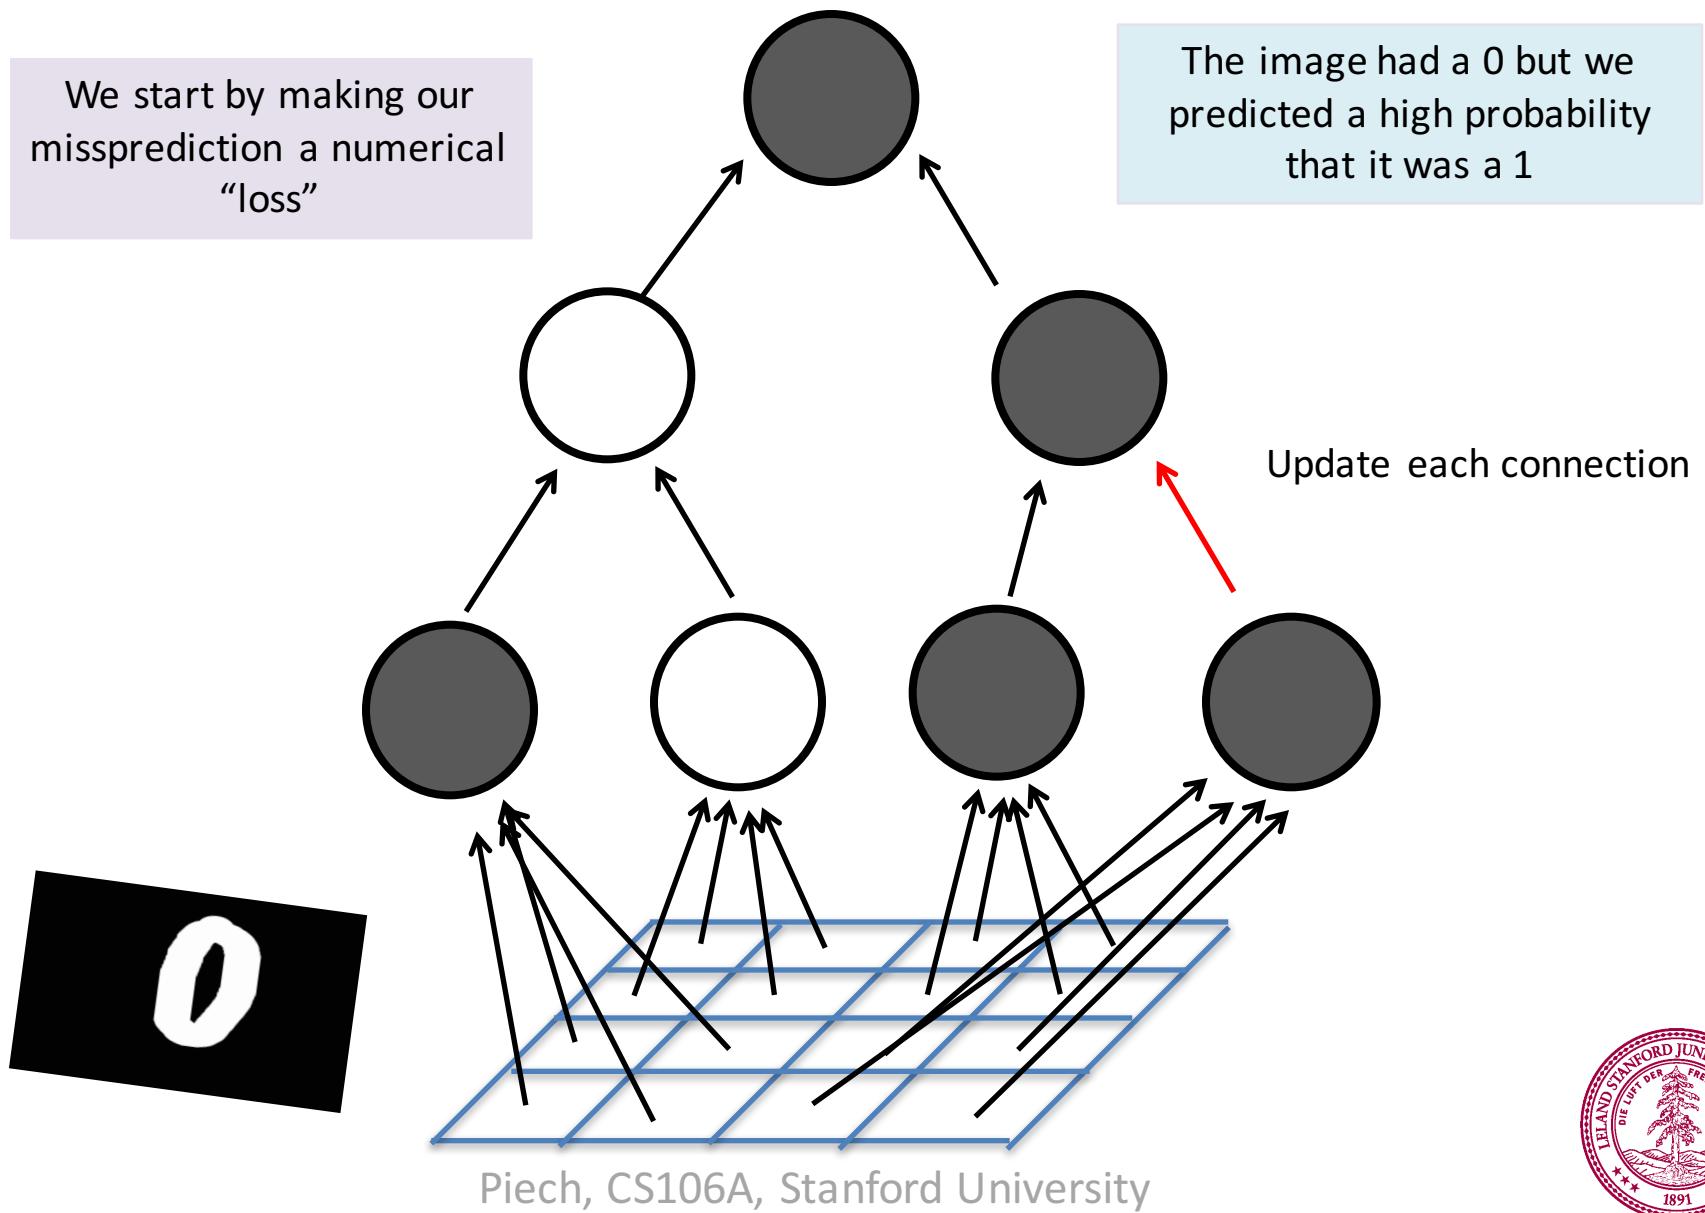
Piech, CS106A, Stanford University



Backward Pass

We start by making our missprediction a numerical “loss”

The image had a 0 but we predicted a high probability that it was a 1



Chose weights that maximize the probability of the right answers

$$P(Y = 1 | X = \mathbf{x}) = \hat{y}$$
$$\hat{y} = \sigma \left(\sum_{j=0}^{m_h} \mathbf{h}_j \theta_j^{(\hat{y})} \right)$$

For one datum

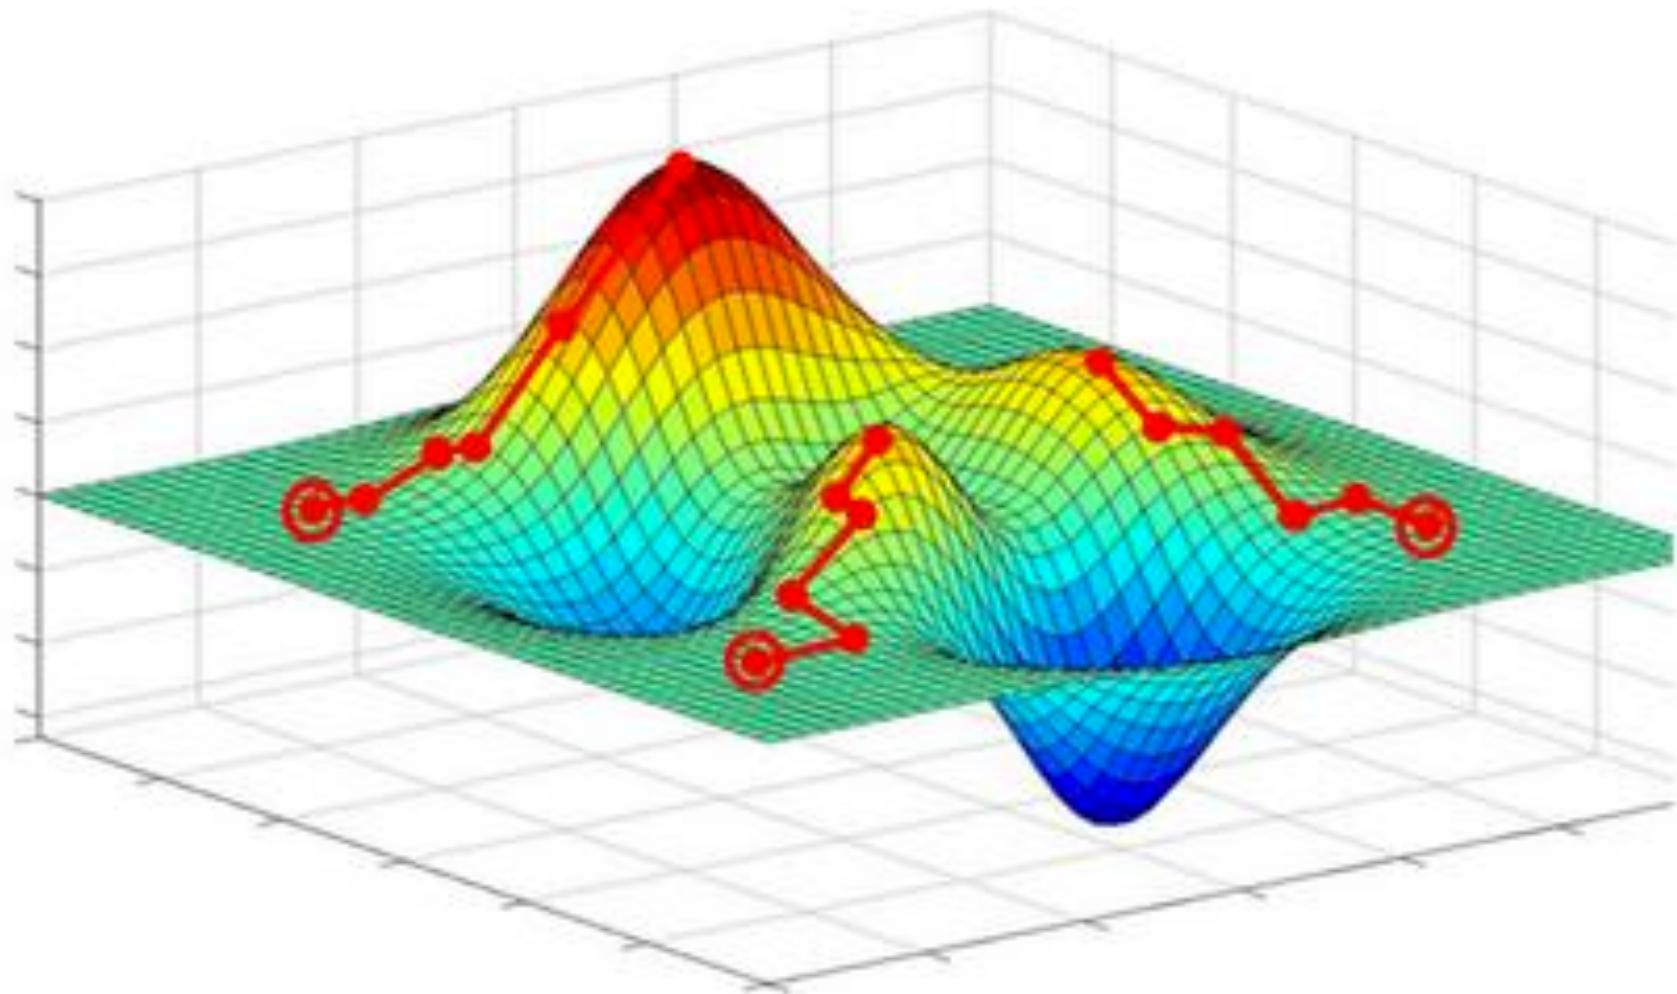
$$P(Y = y | X = \mathbf{X}) = (\hat{y})^y (1 - \hat{y})^{1-y}$$

For IID data

$$L(\theta) = \prod_{i=1}^n P(Y = y^{(i)} | X = \mathbf{x}^{(i)})$$
$$= \prod_{i=1}^n (\hat{y}^{(i)})^{y^{(i)}} \cdot \left[1 - (\hat{y}^{(i)}) \right]^{(1-y^{(i)})}$$



Gradient Ascent



Walk uphill and you will find a local maxima
(if your step size is small enough)
Piech, CS106A, Stanford University



Gradient of output layer params

$$\frac{\partial L}{\partial \theta_i^{(\hat{y})}} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial \theta_i^{(\hat{y})}}$$

$$\hat{y} = \sigma \left(\sum_{j=0}^{m_h} \mathbf{h}_j \theta_j^{(\hat{y})} \right)$$

$$\frac{\partial \hat{y}}{\partial \theta_i^{(\hat{y})}} = \sigma \left(\sum_{j=0}^{m_h} \mathbf{h}_j \theta_j^{(\hat{y})} \right) \left[1 - \sigma \left(\sum_{j=0}^{m_h} \mathbf{h}_j \theta_j^{(\hat{y})} \right) \right] \cdot \frac{\partial}{\partial \theta_i^{(\hat{y})}} \sum_{j=0}^{m_h} \mathbf{h}_j \theta_j^{(\hat{y})}$$

$$= \hat{y}[1 - \hat{y}] \cdot \frac{\partial}{\partial \theta_i^{(\hat{y})}} \sum_{j=0}^{m_h} \mathbf{h}_j \theta_j^{(\hat{y})}$$

$$= \hat{y}[1 - \hat{y}] \cdot h_i$$

That looks scarier than it is



Chain Rule Down the Network

$$\frac{d}{d\theta} = \frac{d}{d\theta} \times \frac{d}{d\theta}$$



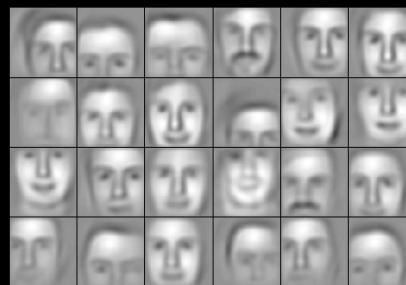
Where you will be by the end of class

When you train,
something really neat happens

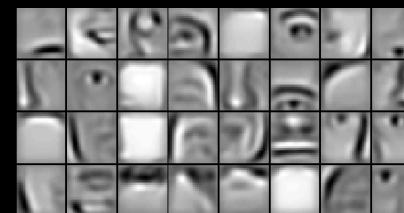
Visualize the Weights



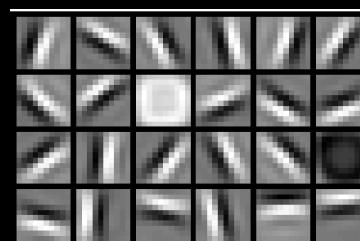
Training set: Aligned
images of faces.



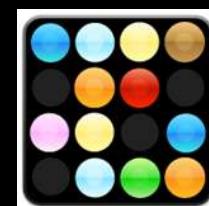
object models



object parts
(combination
of edges)

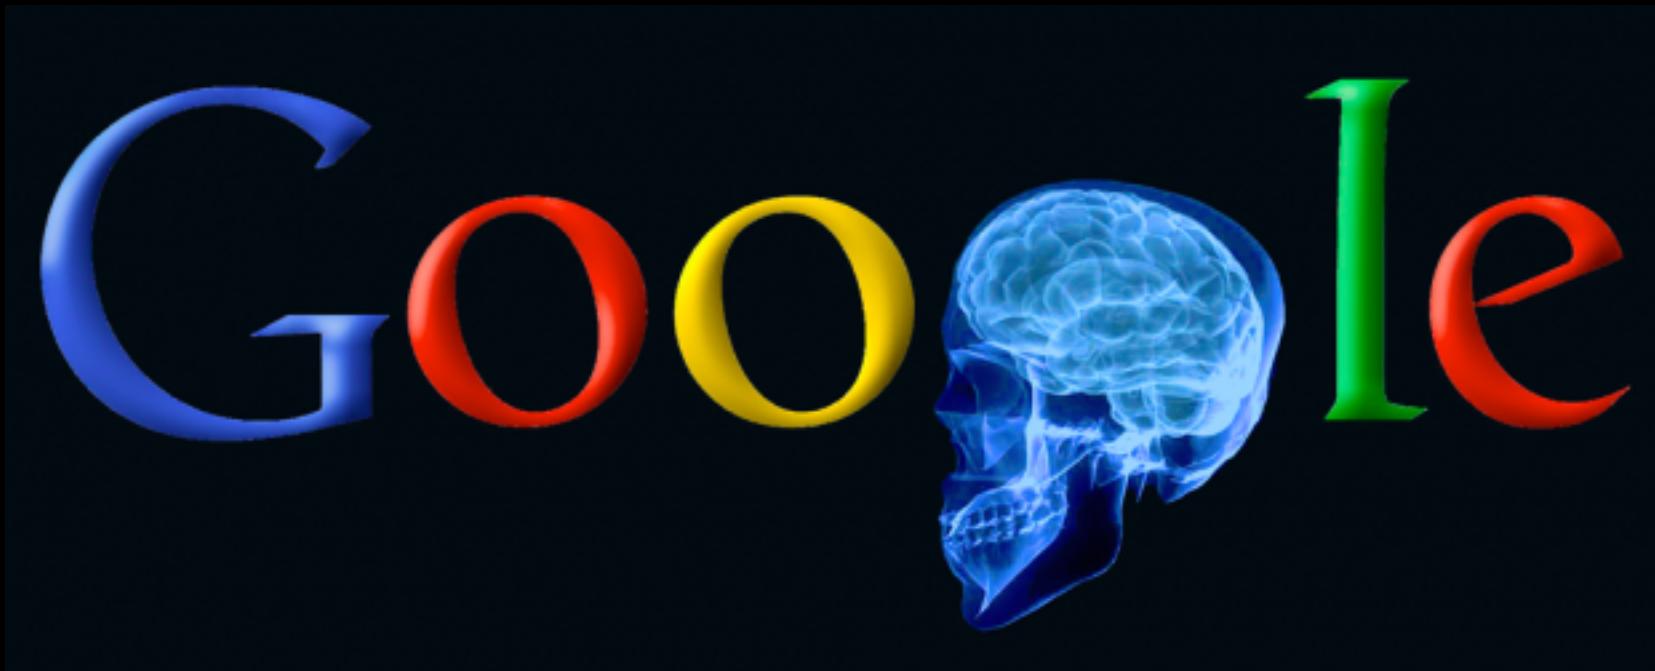


edges

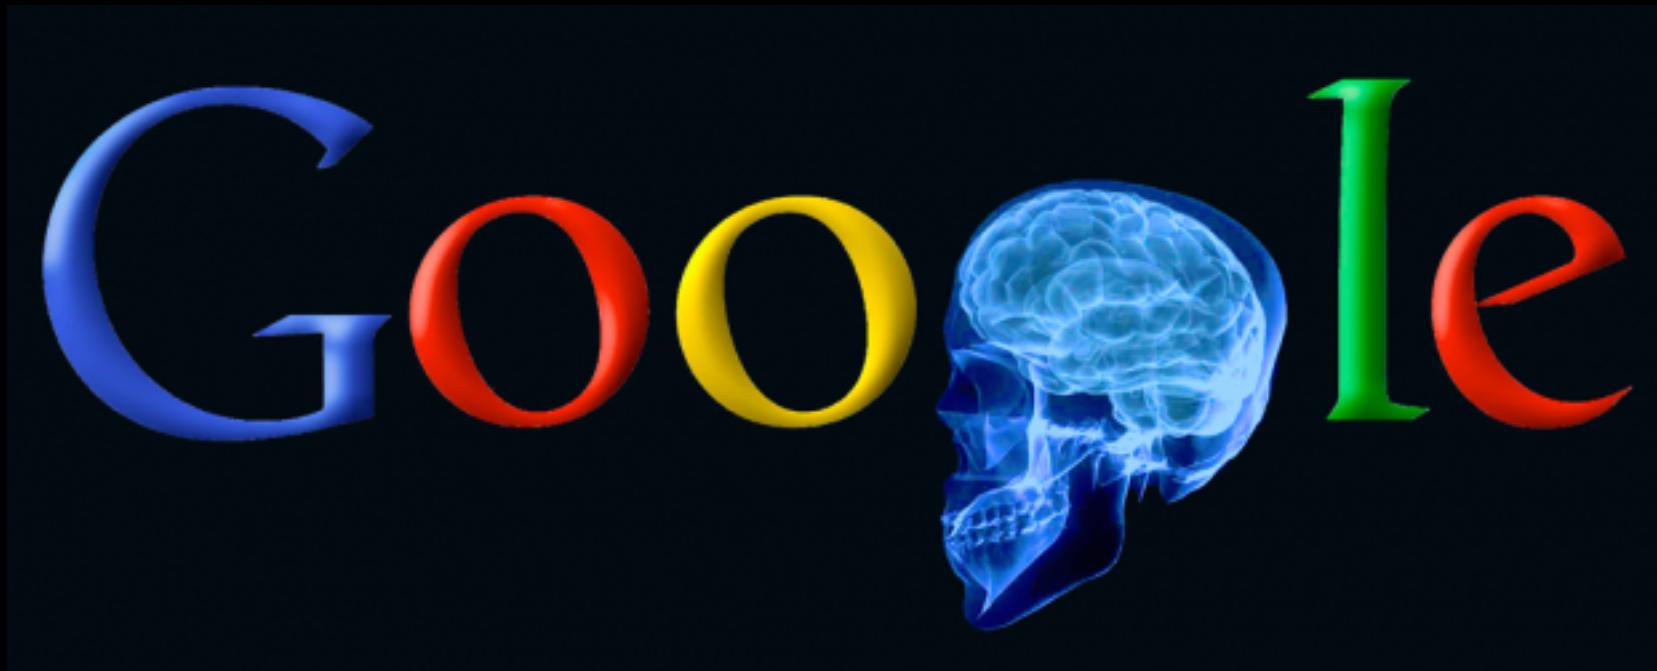


pixels

Google Brain



Google Brain



1 Trillion Artificial Neurons

A Neuron That Fires When It Sees Cats



Top stimuli from the test set



Optimal stimulus
by numerical optimization



We're essentially

a cat detection company

Other Neurons

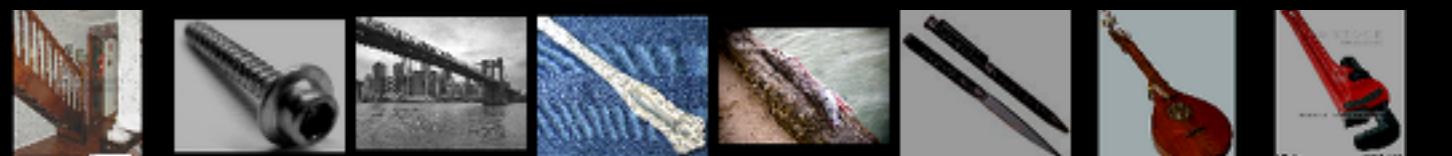
Neuron 1



Neuron 2



Neuron 3



Neuron 4



Neuron 5



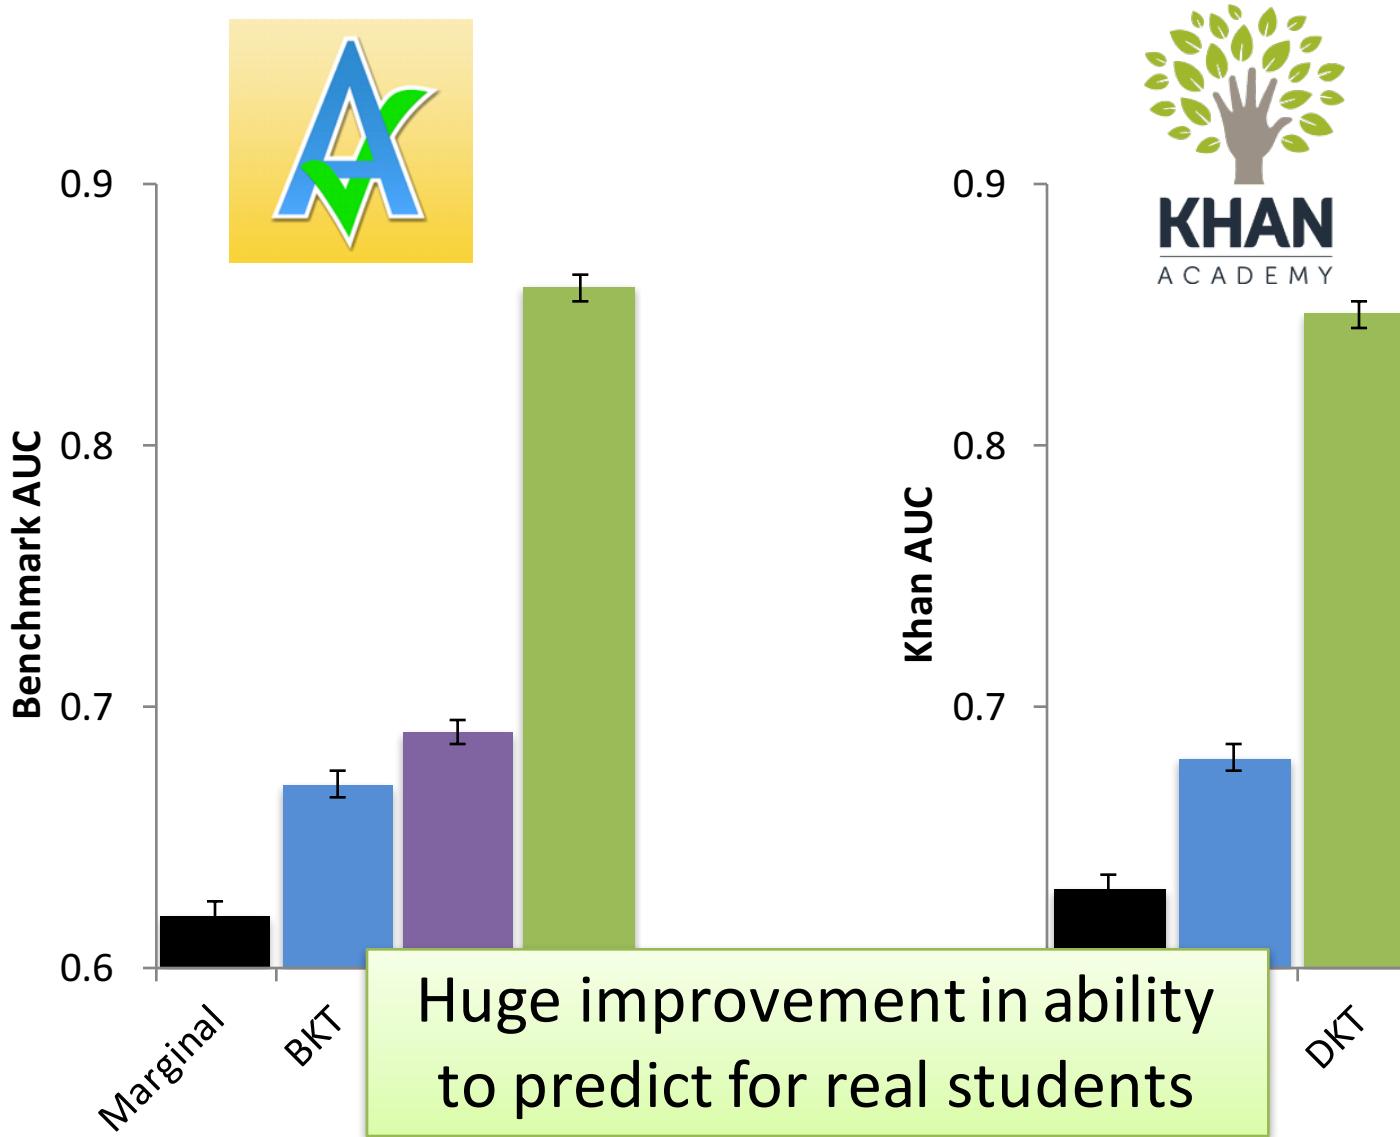
Autonomous Tutor



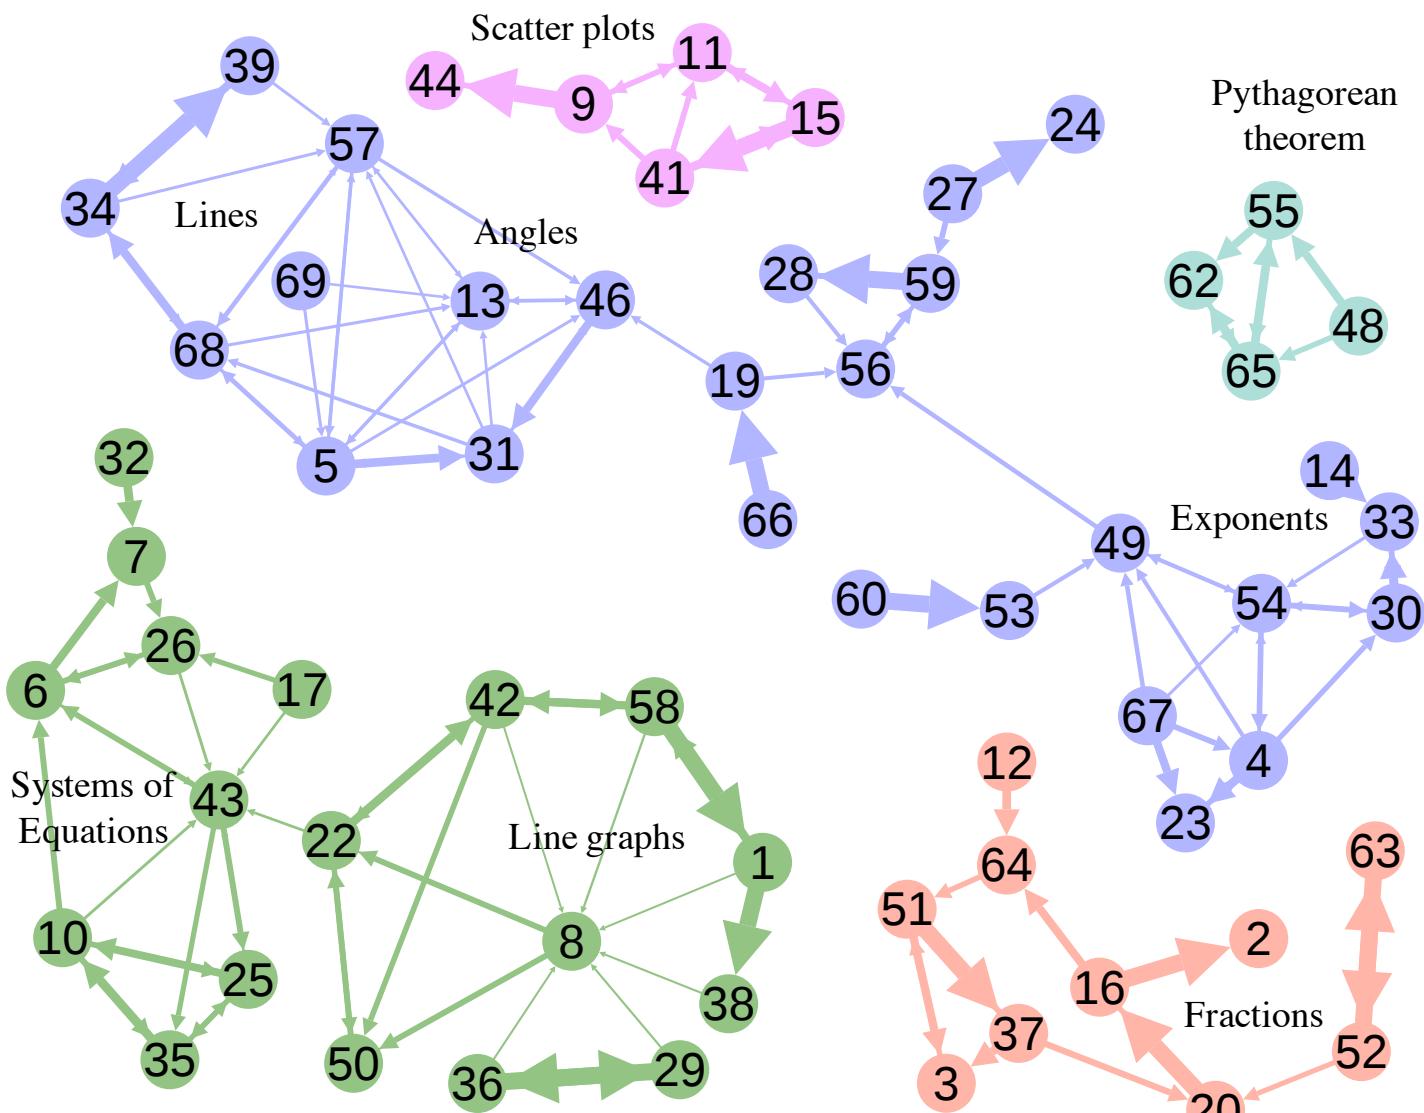
Piech, CS106A, Stanford University



Prediction Results



Autonomous Tutor



Not once, but twice, AI was revolutionized by people who understood probability theory.

End of Story

Except it isn't the end of the story...

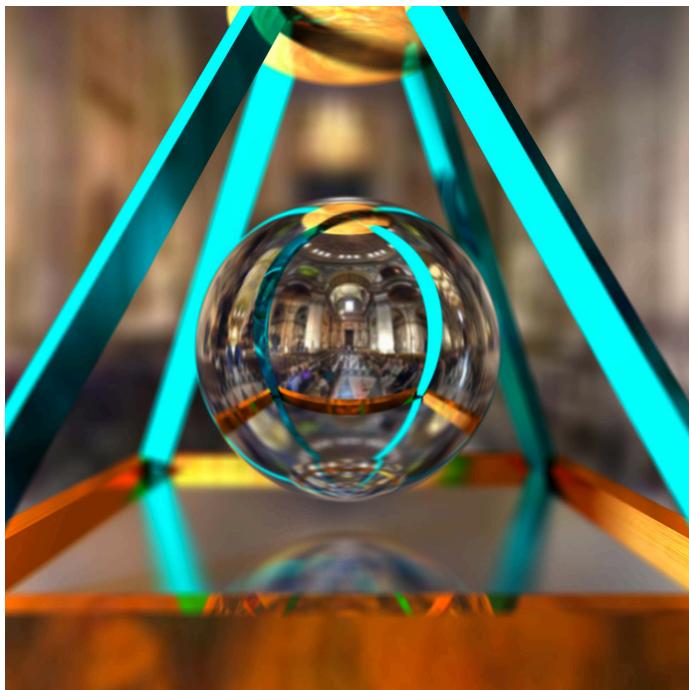
Probability is more than just machine learning

Abundance of Important Problems

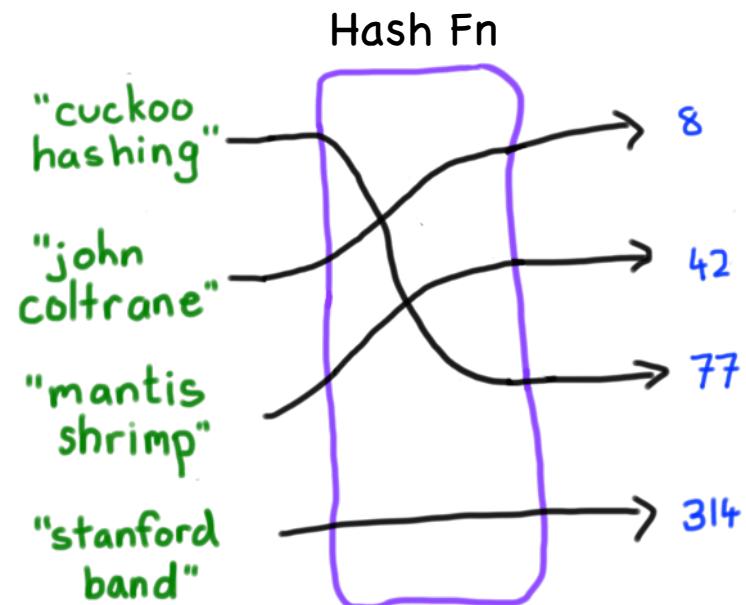


Algorithms and Probability

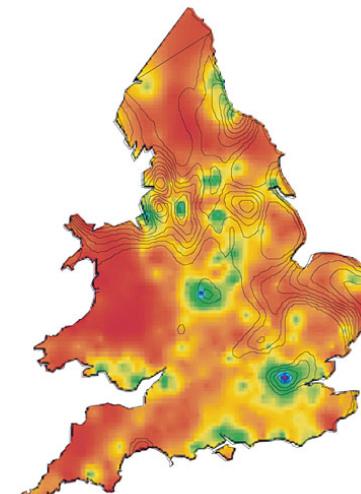
Eg Raytracing



Eg HashMaps



Medicine and Probability



Piech, CS106A, Stanford University

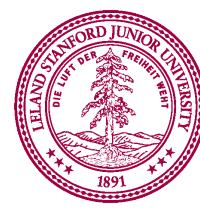


Autocomplete

The image shows a Google search results page. At the top is the Google logo. Below it is a search bar containing the partial query "dinosaurs we". To the right of the search bar are two links: "Advanced Search" and "Language Tools". The main content area displays a list of autocomplete suggestions:

- dinosaurs we
- dinosaurs websites for kids**
- dinosaurs we're back
- dinosaurs webcomic
- dinosaurs webquest
- dinosaurs were made up by the cia to discourage time travel**
- dinosaurs website
- dinosaurs went extinct
- dinosaurs weight
- dinosaurs we are scientists**
- dinosaurs weed episode**

At the bottom of the suggestions are two buttons: "Google Search" and "I'm Feeling Lucky".



Probability in Practice

amazon.com

Hello. Sign in to get personalized recommendations. New customer? Start here.

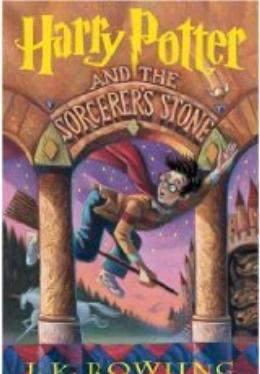
FREE 2-Day Shipping, No Minimum Purchase

Your Amazon.com | Today's Deals | Gifts & Wish Lists | Gift Cards

Your Account | Help

Shop All Departments | Search Books | GO | Cart | Your Lists

Books | Advanced Search | Browse Subjects | Hot New Releases | Bestsellers | The New York Times® Best Sellers | Libros En Español | Bargain Books | Textbooks



Harry Potter and the Sorcerer's Stone (Book 1) (Hardcover)
by J.K. Rowling (Author), Mary GrandPré (Illustrator)
 (5,471 customer reviews)

List Price: \$24.99
Price: **\$15.92** & eligible for **FREE Super Saver Shipping** on orders over \$25.
[Details](#)
You Save: **\$9.07 (36%)**

In Stock.
Ships from and sold by **Amazon.com**. Gift-wrap available.

Quantity: 1
or
[Sign in](#) to turn on 1-Click ordering.
or

Amazon Prime Free Trial required. Sign up when you check out. [Learn More](#)

Customers Who Bought This Item Also Bought


[Harry Potter and the Prisoner of Azkaban \(Book 3\) by J.K. Rowling](#)
★★★★★ (2,599) \$16.49


[Harry Potter and the Goblet of Fire \(Book 4\) by J.K. Rowling](#)
★★★★★ (5,186) \$19.79


[Harry Potter and the Order of the Phoenix \(Book 5\) by J. K. Rowling](#)
★★★★★ (5,876) \$10.18

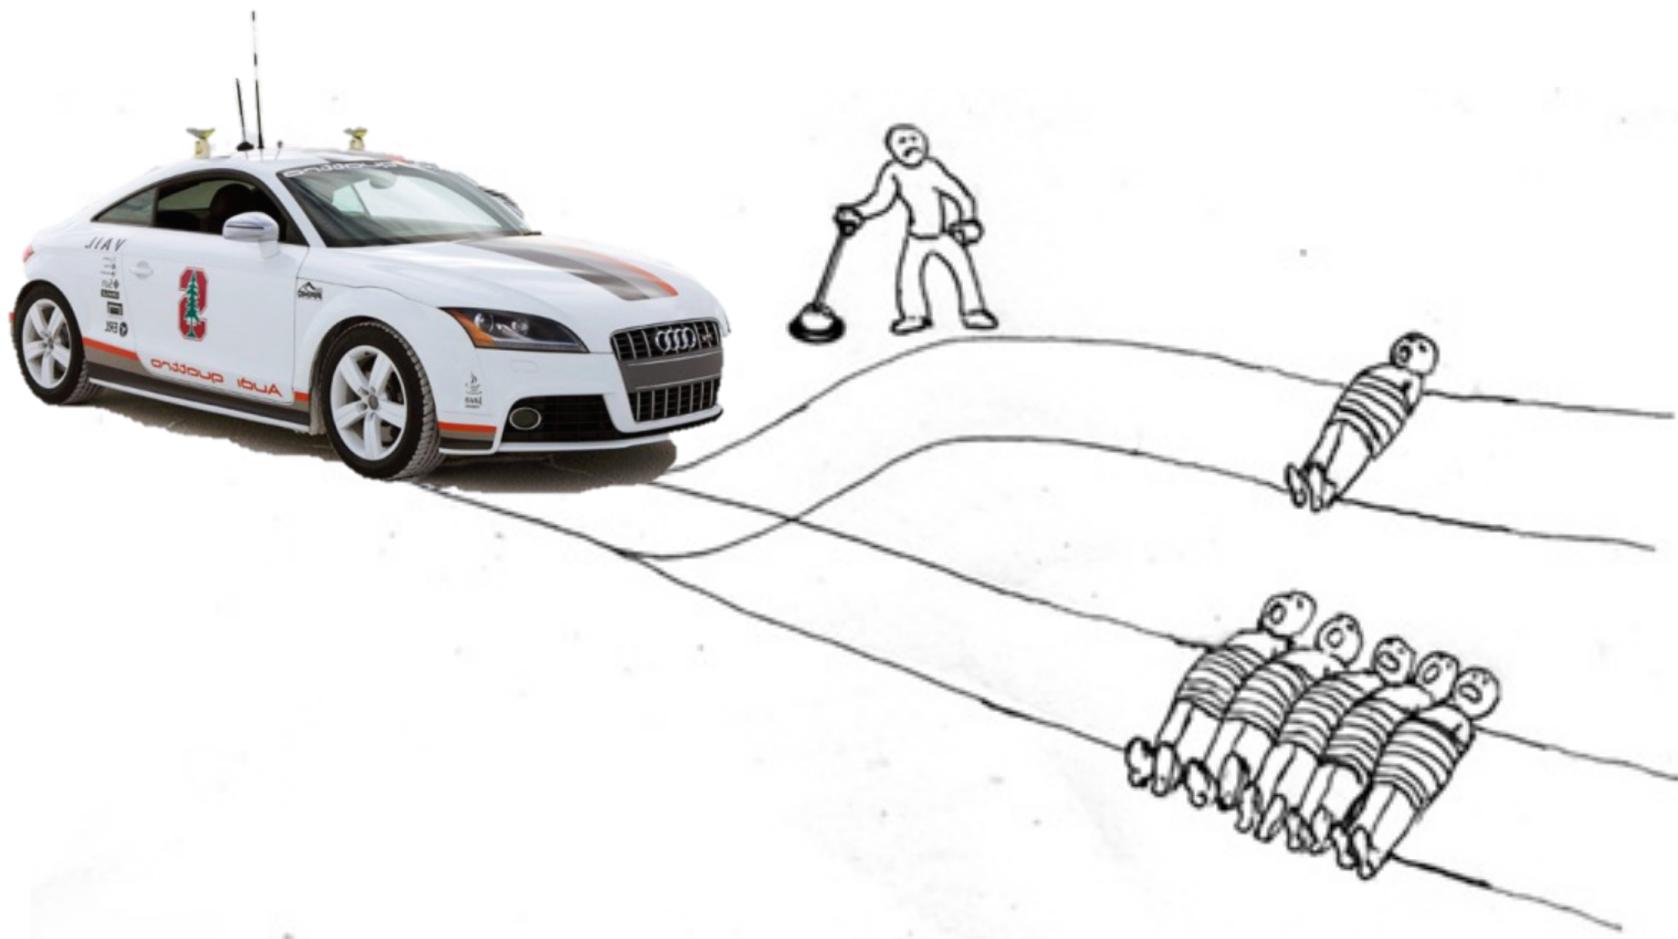

[Harry Potter and the Half-Blood Prince \(Book 6\) by J.K. Rowling](#)
★★★★★ (3,597) \$10.18


[The Tales of Beedle the Bard, Collector's Ed... by J. K. Rowling](#)
★★★★★ (176)

Page 1 of 20



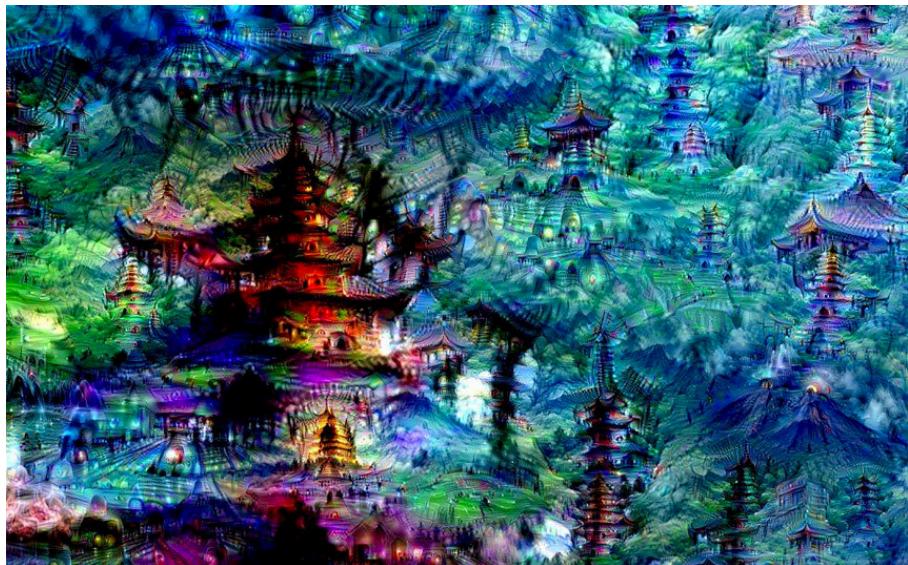
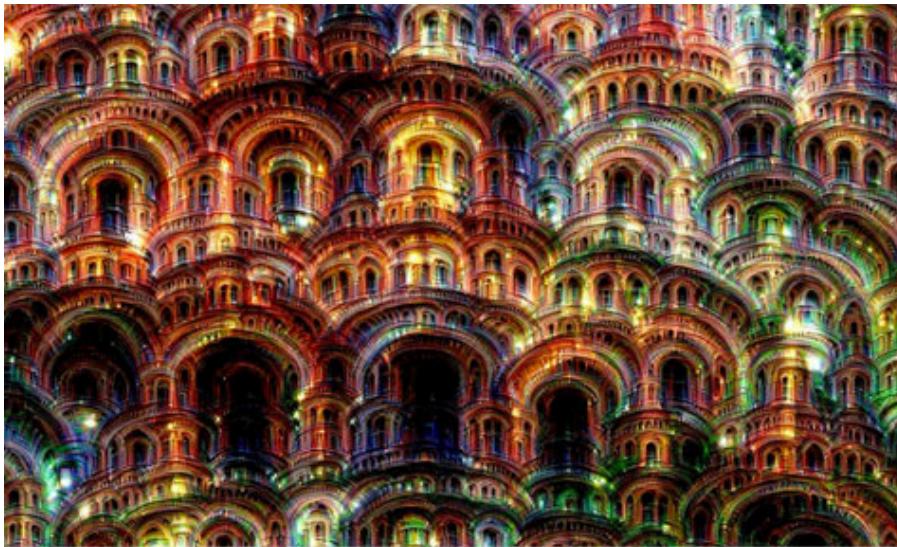
Philosophy and Probability



Piech, CS106A, Stanford University



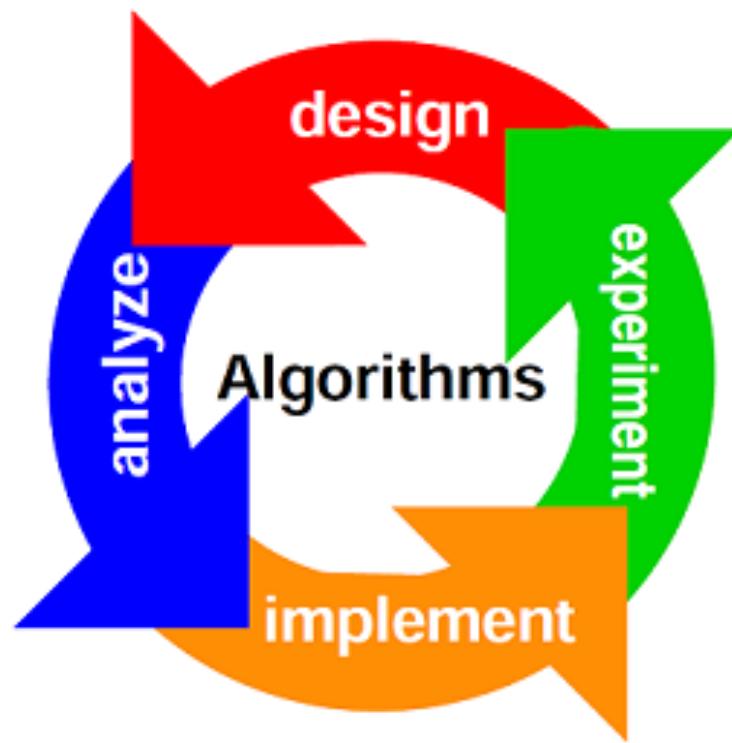
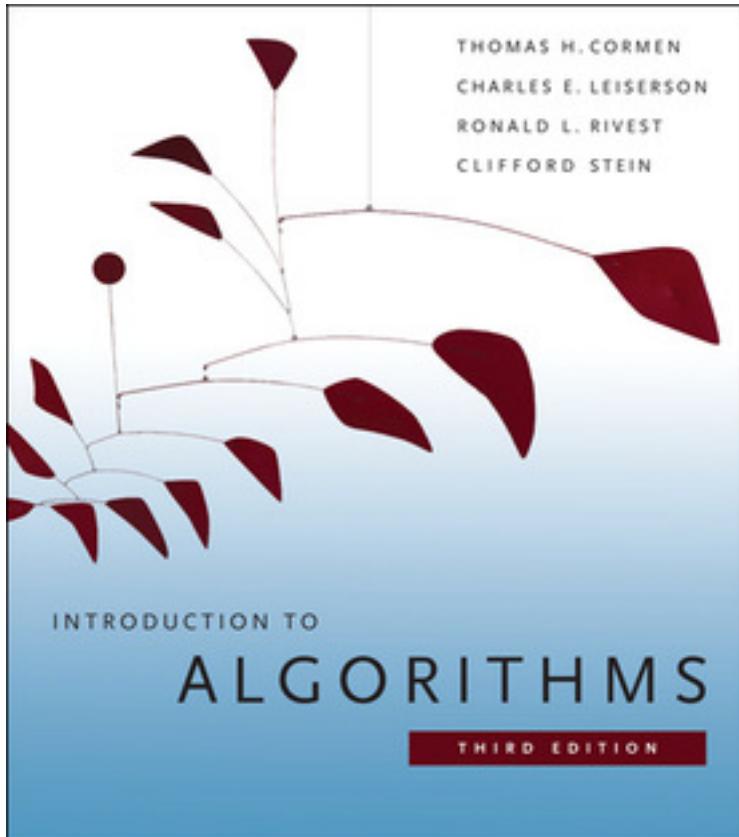
Art and Probability



Piech, CS106A, Stanford University



Probabilistic Analysis of Algorithms



#1 Most Desired Skill in Industry

Microsoft's competitive advantage, [Bill Gates] responded, was its expertise in "Bayesian [probabilistic] networks."

(from Los Angeles Times, Oct. 28, 1996)

"The sexy job in the next 10 years will be statisticians."

-Hal Varian, Chief Economist at Google
(from New York Times, August 6, 2009)



#1 Most Desired Skill in Industry

“I believe over the next decade computing will become even more ubiquitous and intelligence will become ambient. The coevolution of software and new hardware form factors will intermediate and digitize — many of the things we do and experience in business, life and our world. This will be made possible by an ever-growing network of connected devices, incredible computing capacity from the cloud, insights from big data, and intelligence from machine learning.”

-- Satya Nadella (CEO, Microsoft)

Email to all employees on first day as CEO (Feb. 04, 2014)



#1 Most Desired Skill in Academia

Most CS PhD students list their highest desiderata upon graduation as:

“Better understanding of probability”



Foundation for your future

But its not always intuitive

Zika Test



Positive Zika.

What is the probability of zika?

- *0.08% of people have zika*
- *90% positive rate for people with zika*
- *7% positive rate for people without zika*

The right answer is 9%

Piech, CS106A, Stanford University



Probability = Important + Needs Study

Delayed gratification

What is CS109?

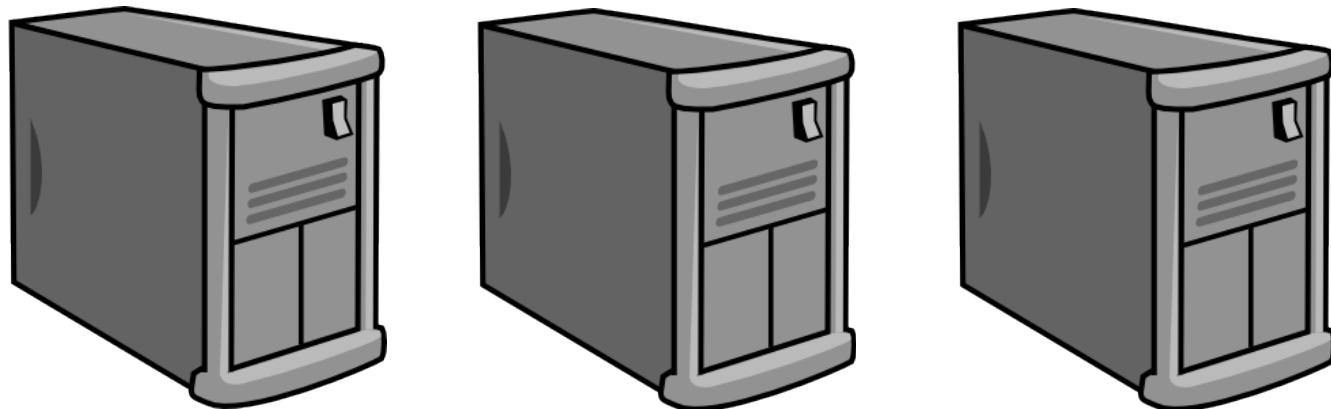
Traditional View of Probability



CS View of Probability

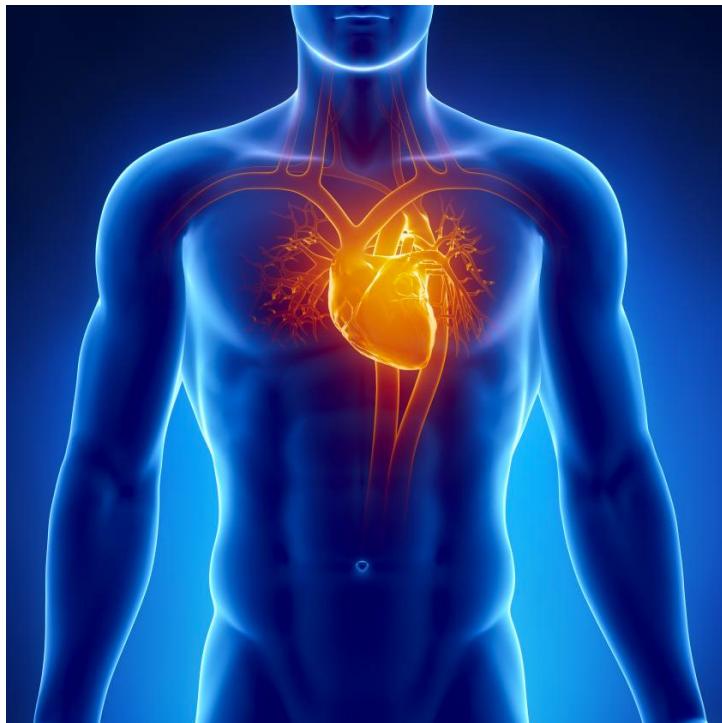
<http://www.site.com>

Give you the tools necessary to build and understand probabilistic CS algorithms.



CS View of Probability

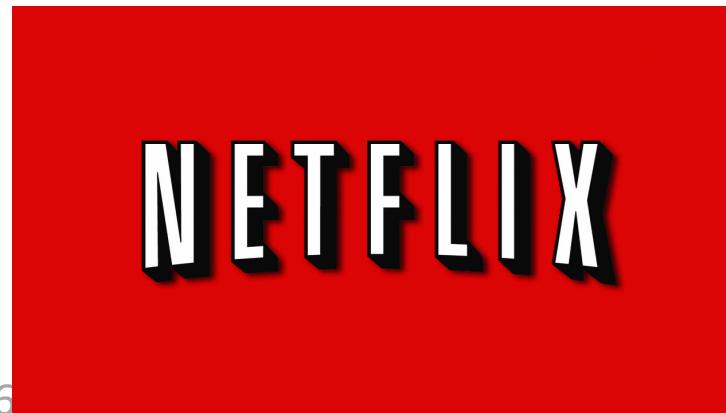
Heart



Ancestry



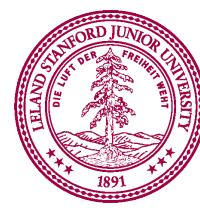
Netflix



CS View of Probability



Piech, CS106A, Stanford University



CS View of Probability

Teach you how to write programs
that most people are not able to write.

Spring vs Winter?



Lets dive in...

Counting



Our Route

